

Sequence Generation with Recurrent Neural Networks

Arne Nix

arne.nix@rwth-aachen.de

**Selected Topics in Human Language Technology and Pattern Recognition
WS 2015/2016 – 19.01.2016**

**Human Language Technology and Pattern Recognition
Lehrstuhl für Informatik 6
Computer Science Department
RWTH Aachen University, Germany**

Literature

A. Graves [Graves 13]:

Generating sequences with recurrent neural networks. *August 2013.*

- ▶ Basic concept of generating sequences
- ▶ Application to text prediction and handwriting prediction

I. Sutskever, O. Vinyals, Q. Le [Sutskever & Vinyals⁺ 14]:

Sequence to sequence learning with neural networks. *NIPS December 2014.*

- ▶ Generating sequences from sequences using recurrent neural networks.
- ▶ Introducing the encoder-decoder model
- ▶ Application to machine translation

O. Vinyals, Q. Le [Vinyals & Le 15]:

A neural conversational model. *ICML July 2015.*

- ▶ Using the RNN encoder-decoder model to model conversations

Outline

- 1. Literature**
- 2. Introduction**
- 3. Recap: Recurrent Neural Networks**
- 4. Generating Sequences**
- 5. Conversation Modelling**
 - (a) Sequence to Sequence Approach**
 - (b) Alternative Approaches**
 - (c) Evaluation**
- 6. Conclusion**

Introduction

We want to solve the problem of *sequence to sequence generation*

► **Training set:**

- **Source sequences** $\bar{x} = (x(1), \dots, x(T)) \in (\mathbb{R}^n \times \dots \times \mathbb{R}^n)$
- **Target sequences** $\bar{y} = (y(1), \dots, y(T')) \in (\mathbb{R}^n \times \dots \times \mathbb{R}^n)$

► **Goal:**

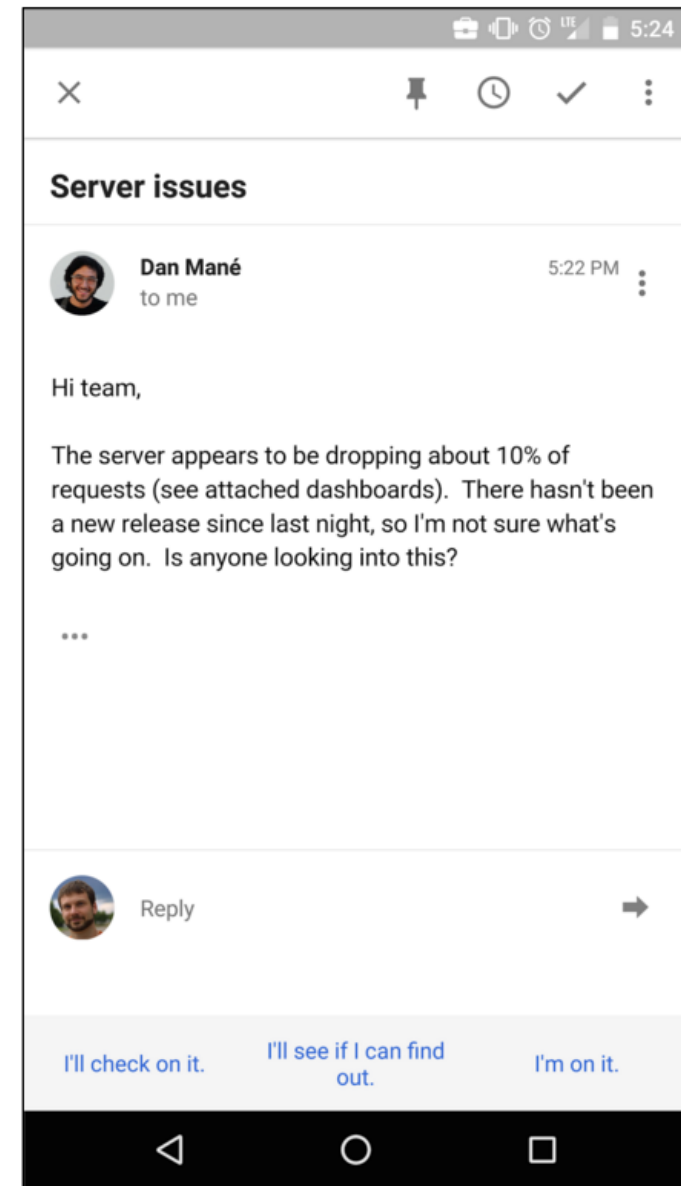
- **Generate output sequence** \bar{y}^* for unseen test sequence \bar{x}^*
- **Maximize** $\Pr(\bar{y}^* | \bar{x}^*)$

Motivation

Possible motivations:


- ▶ Simulate future situations
- ▶ Achieve a better representation of the data
- ▶ Practical applications:
 - ▷ *Speech synthesis*
 - ▷ *Machine translation*
 - ▷ *Response generation*
 - ▷ *Handwriting generation*
 - ▷ *Conversation modelling*

handwriting generation



Google Inbox

Neural Networks



Feed forward neural network

► **Activation of neuron i :**

$$y_i = f_i\left(\sum_{j=1}^J w_{ij} \cdot x_j + b_i\right) \quad (1)$$

► **Activation vector for layer l :**

$$y^{(l)} = f^{(l)}(W_{\text{in}}^{(l)} y^{(l-1)} + b^{(l)}) \quad (2)$$

► **Common activation functions:**


$$\sigma(x) = \frac{1}{1 + e^{-x}} \quad (3)$$

$$\tanh(x) = \frac{e^{2x} - 1}{e^{2x} + 1} \quad (4)$$

Recurrent Neural Networks

► Activation of layer l for timestep t :

$$y^{(l)}(t) = f^{(l)}(W_{\text{in}}^{(l)} y^{(l-1)}(t) + W_{\text{re}}^{(l)} y^{(l)}(t-1) + b^{(l)}) \quad (5)$$



RNN with its equivalent unfolded in time for three time steps.

Training Neural Networks with Gradient Descent

- Small step of fixed size α in the direction of the negative error gradient


$$\Delta w_{ji}^{(l),n} = -\alpha \frac{\partial \mathcal{L}}{\partial w_{ji}^{(l),n}} \quad (6)$$

where $\Delta w_{ji}^{(l),n}$ is the update of $w_{ji}^{(l)}$ for the n^{th} iteration and loss function \mathcal{L}

- Momentum parameter $m \in [0, 1]$ for faster convergence

$$\Delta w_{ji}^{(l),n} = m \Delta w_{ji}^{(l),n-1} - \alpha \frac{\partial \mathcal{L}}{\partial w_{ji}^{(l),n}} \quad (7)$$

- Gradient calculated by *backpropagation* or *backpropagation through time*



Training Neural Networks with Gradient Descent

- Small step of fixed size α in the direction of the negative error gradient


$$\Delta w_{ji}^{(l),n} = -\alpha \frac{\partial \mathcal{L}}{\partial w_{ji}^{(l),n}} \quad (6)$$

where $\Delta w_{ji}^{(l),n}$ is the update of $w_{ji}^{(l)}$ for the n^{th} iteration and loss function \mathcal{L}

- Momentum parameter $m \in [0, 1]$ for faster convergence

$$\Delta w_{ji}^{(l),n} = m \Delta w_{ji}^{(l),n-1} - \alpha \frac{\partial \mathcal{L}}{\partial w_{ji}^{(l),n}} \quad (7)$$

- Gradient calculated by *backpropagation* or *backpropagation through time*



Training Neural Networks with Gradient Descent

- Small step of fixed size α in the direction of the negative error gradient


$$\Delta w_{ji}^{(l),n} = -\alpha \frac{\partial \mathcal{L}}{\partial w_{ji}^{(l),n}} \quad (6)$$

where $\Delta w_{ji}^{(l),n}$ is the update of $w_{ji}^{(l)}$ for the n^{th} iteration and loss function \mathcal{L}

- Momentum parameter $m \in [0, 1]$ for faster convergence

$$\Delta w_{ji}^{(l),n} = m \Delta w_{ji}^{(l),n-1} - \alpha \frac{\partial \mathcal{L}}{\partial w_{ji}^{(l),n}} \quad (7)$$

- Gradient calculated by *backpropagation* or *backpropagation through time*



Training Neural Networks with Gradient Descent

- Small step of fixed size α in the direction of the negative error gradient


$$\Delta w_{ji}^{(l),n} = -\alpha \frac{\partial \mathcal{L}}{\partial w_{ji}^{(l),n}} \quad (6)$$

where $\Delta w_{ji}^{(l),n}$ is the update of $w_{ji}^{(l)}$ for the n^{th} iteration and loss function \mathcal{L}

- Momentum parameter $m \in [0, 1]$ for faster convergence

$$\Delta w_{ji}^{(l),n} = m \Delta w_{ji}^{(l),n-1} - \alpha \frac{\partial \mathcal{L}}{\partial w_{ji}^{(l),n}} \quad (7)$$

- Gradient calculated by *backpropagation* or *backpropagation through time*



Training Neural Networks with Gradient Descent

- Small step of fixed size α in the direction of the negative error gradient


$$\Delta w_{ji}^{(l),n} = -\alpha \frac{\partial \mathcal{L}}{\partial w_{ji}^{(l),n}} \quad (6)$$

where $\Delta w_{ji}^{(l),n}$ is the update of $w_{ji}^{(l)}$ for the n^{th} iteration and loss function \mathcal{L}


- Momentum parameter $m \in [0, 1]$ for faster convergence

$$\Delta w_{ji}^{(l),n} = m \Delta w_{ji}^{(l),n-1} - \alpha \frac{\partial \mathcal{L}}{\partial w_{ji}^{(l),n}} \quad (7)$$

- Gradient calculated by *backpropagation* or *backpropagation through time*



Long Short Term Memory (LSTM)



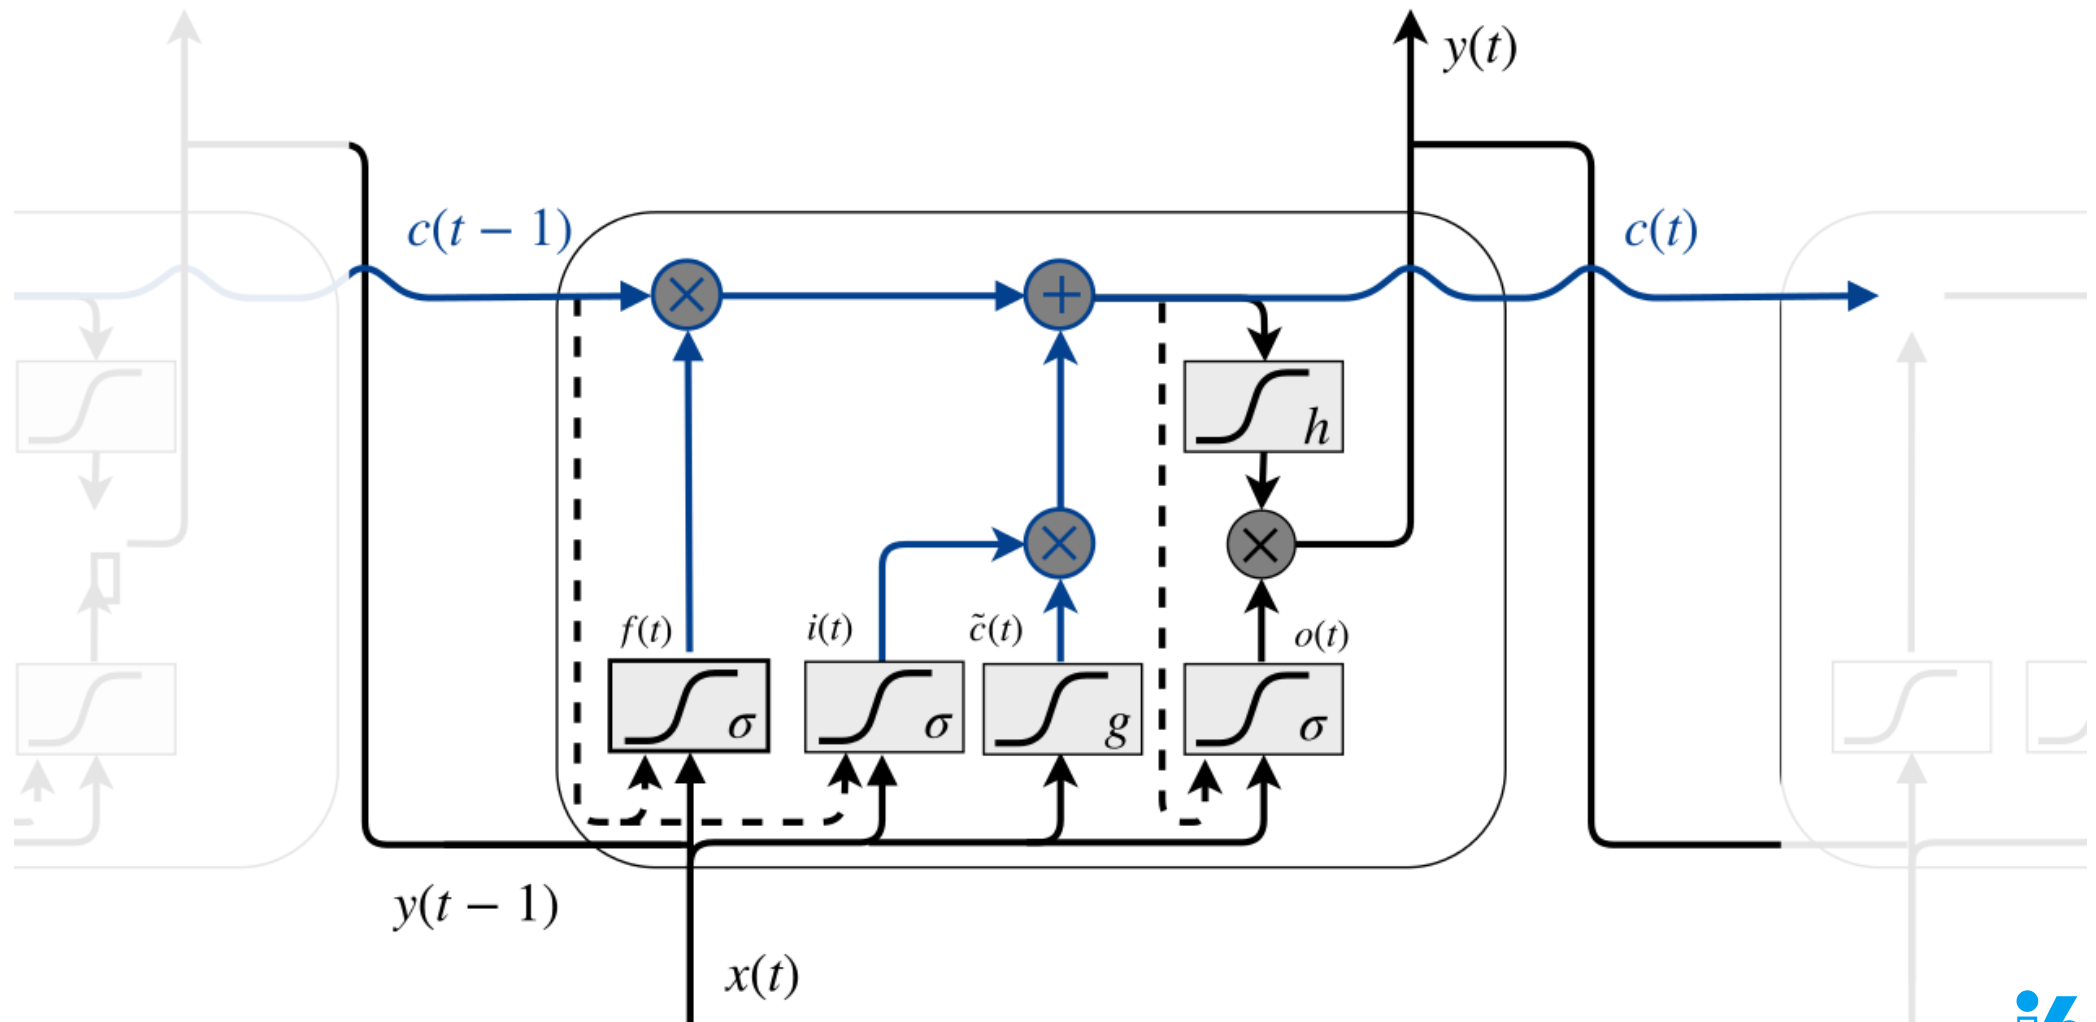
Problem: vanishing gradient

- ▶ Multiplying gradients in the range of $(0, 1]$ or $(0, 0.25]$
- ▶ For timesteps lying further in the past: $\frac{\partial \mathcal{L}}{\partial w_{ji}^{(l)}} \rightarrow 0$
- ▶ Also exploding gradient possible: $\frac{\partial \mathcal{L}}{\partial w_{ji}^{(l)}} \rightarrow \pm \infty$
- ▶ Consequences:
Standard RNNs do not provide long range contextual information
- ▶ Solution:
Long Short Term Memory [Hochreiter & Schmidhuber 97] and variations

Long Short Term Memory (LSTM)

Updated state:

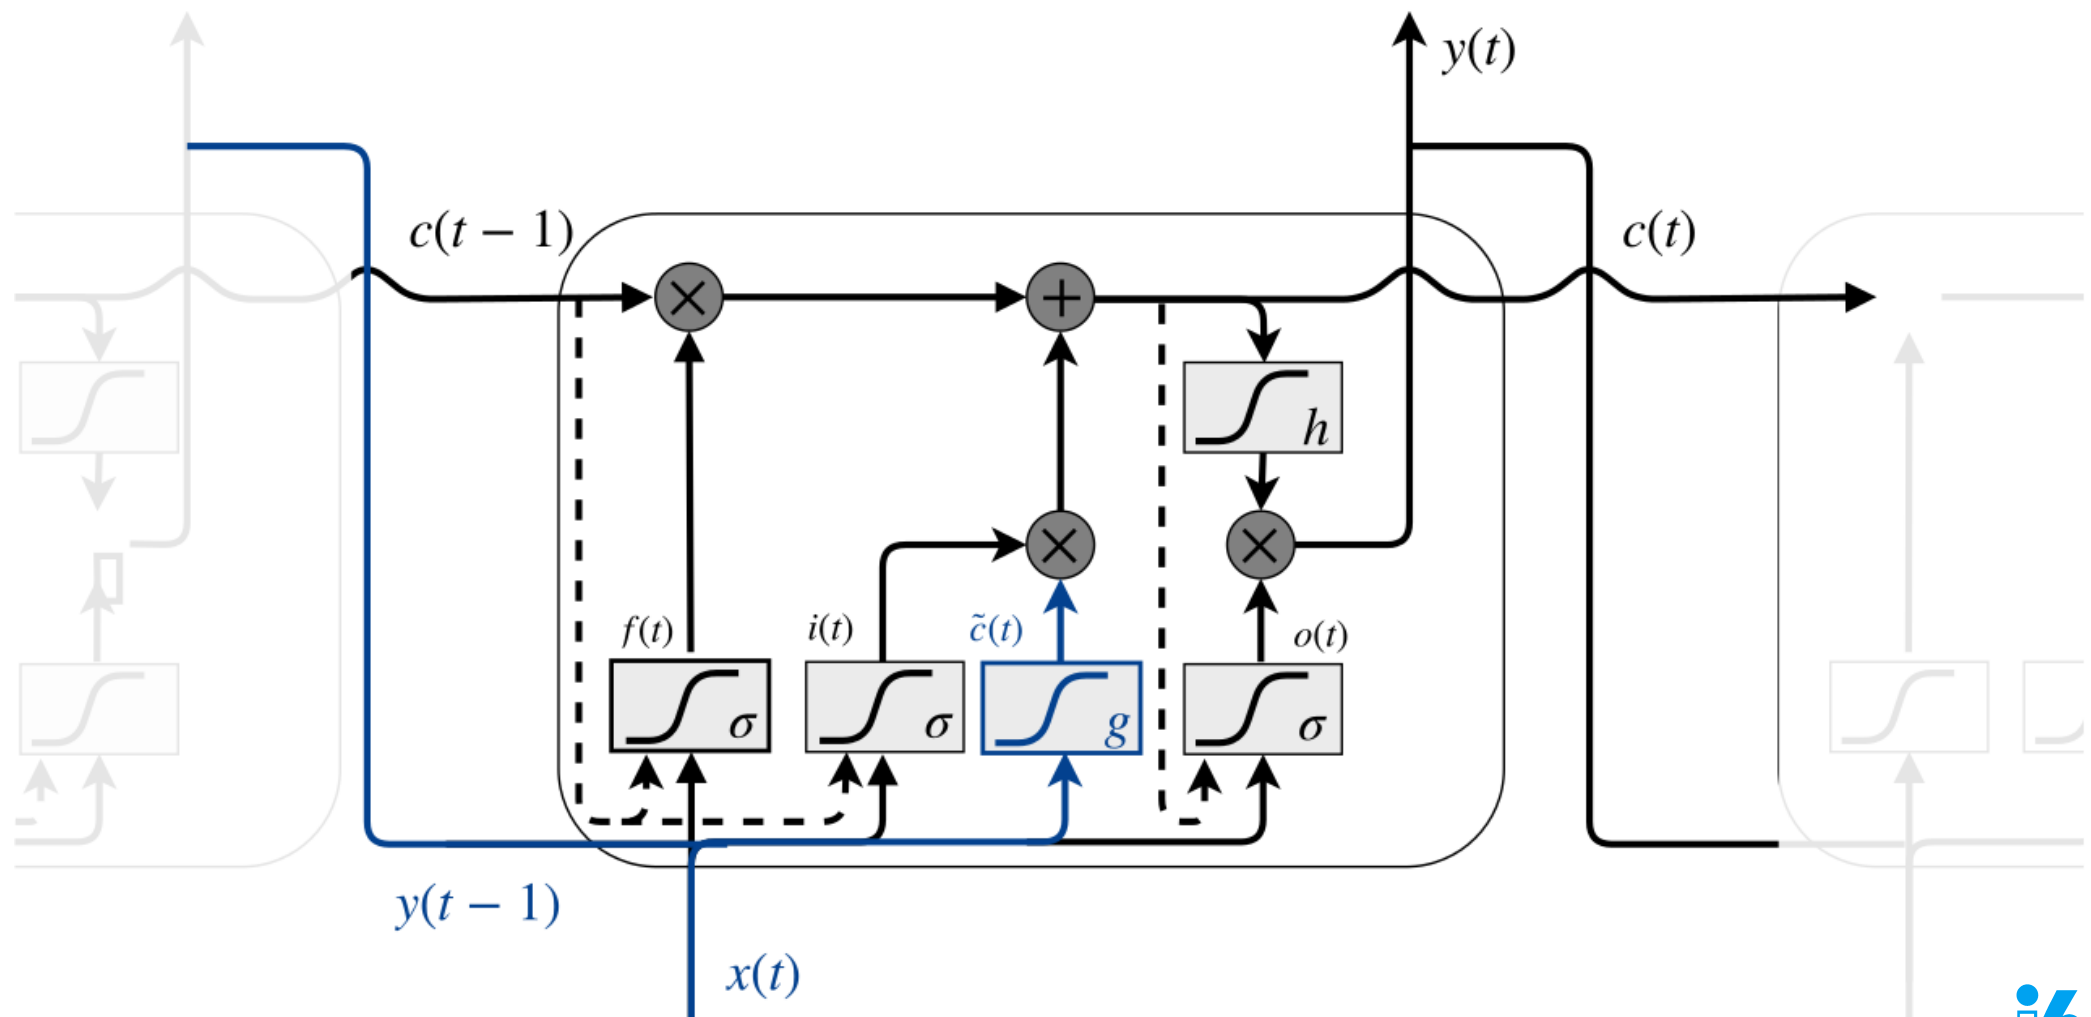
$$c(t) = f(t) \cdot c(t - 1) + i(t) \cdot \tilde{c}(t) \quad (8)$$



Long Short Term Memory (LSTM)

Update candidate

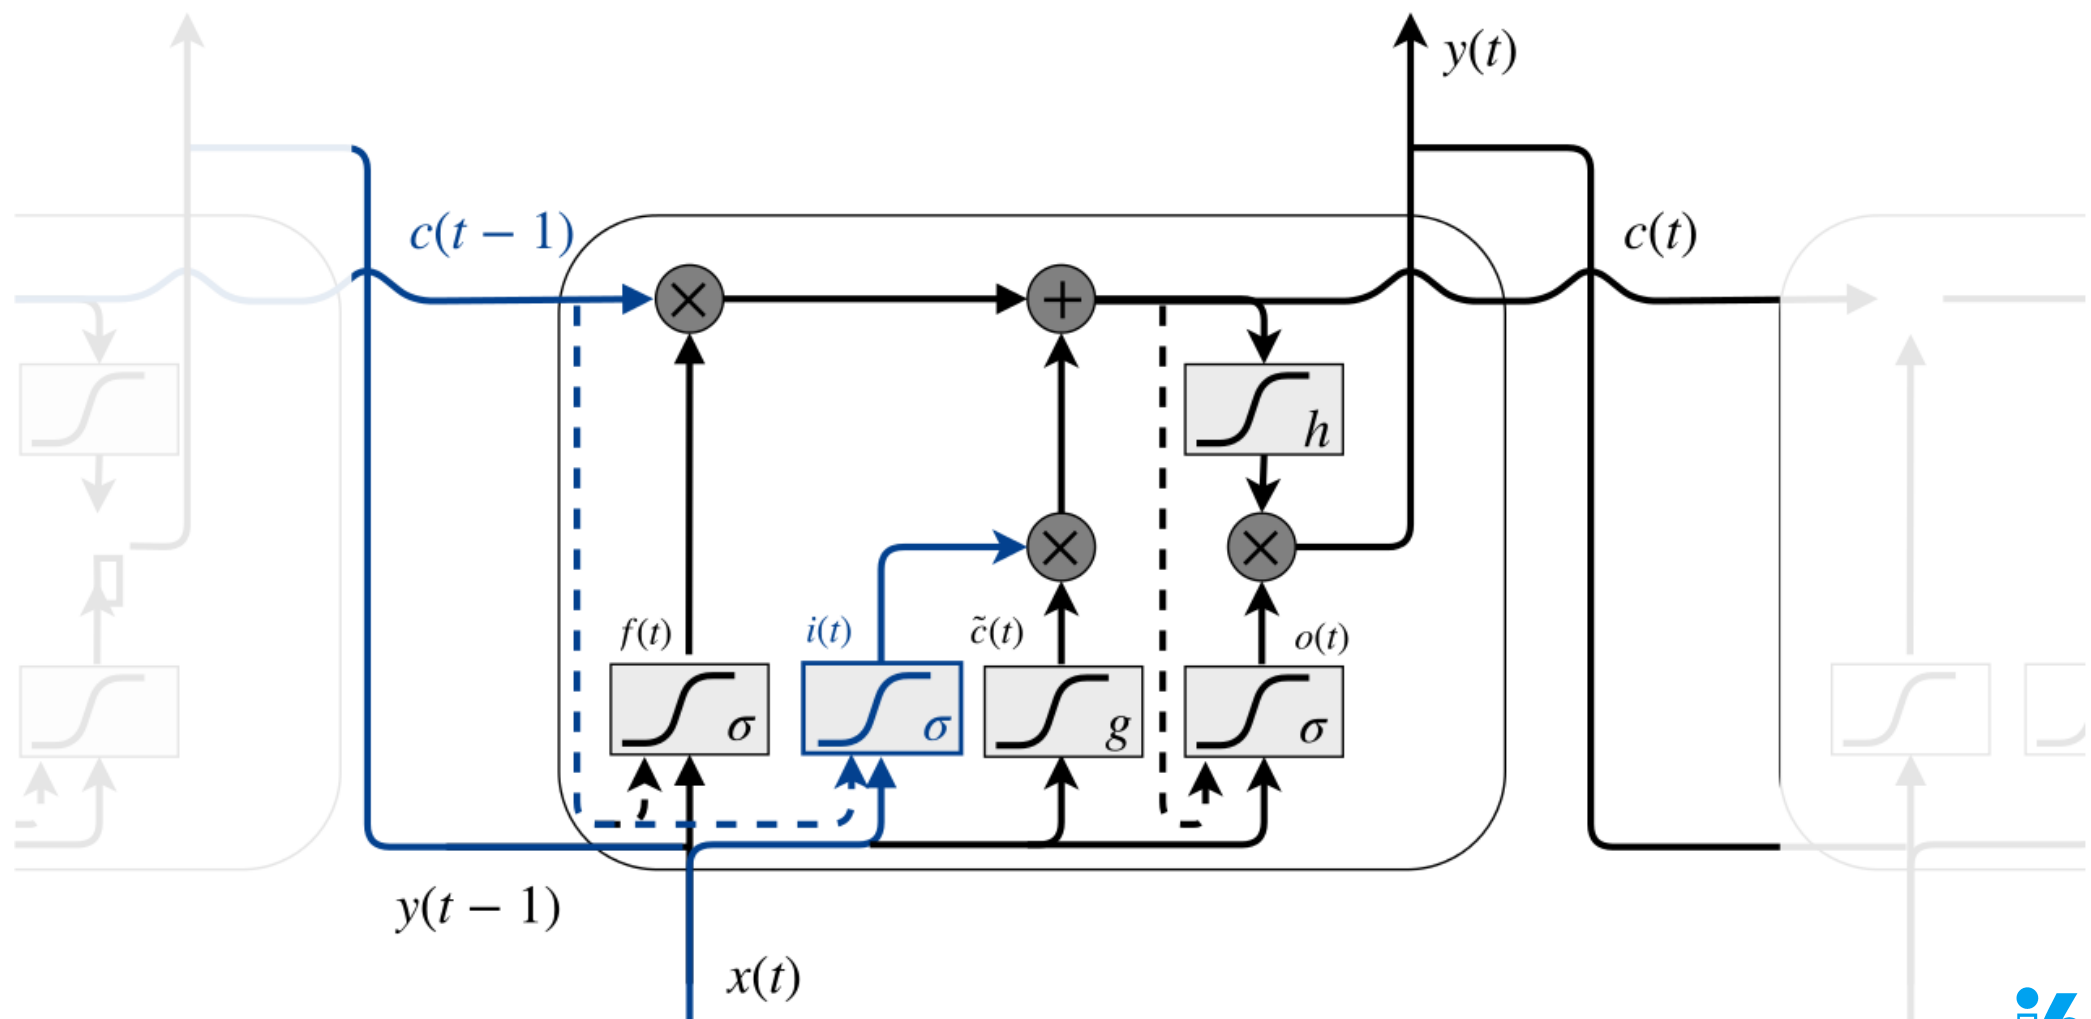
$$\tilde{c}(t) = g(W_{xc}x(t) + W_{yc}y(t-1) + b_c) \quad (9)$$



Long Short Term Memory (LSTM)

Input gate:

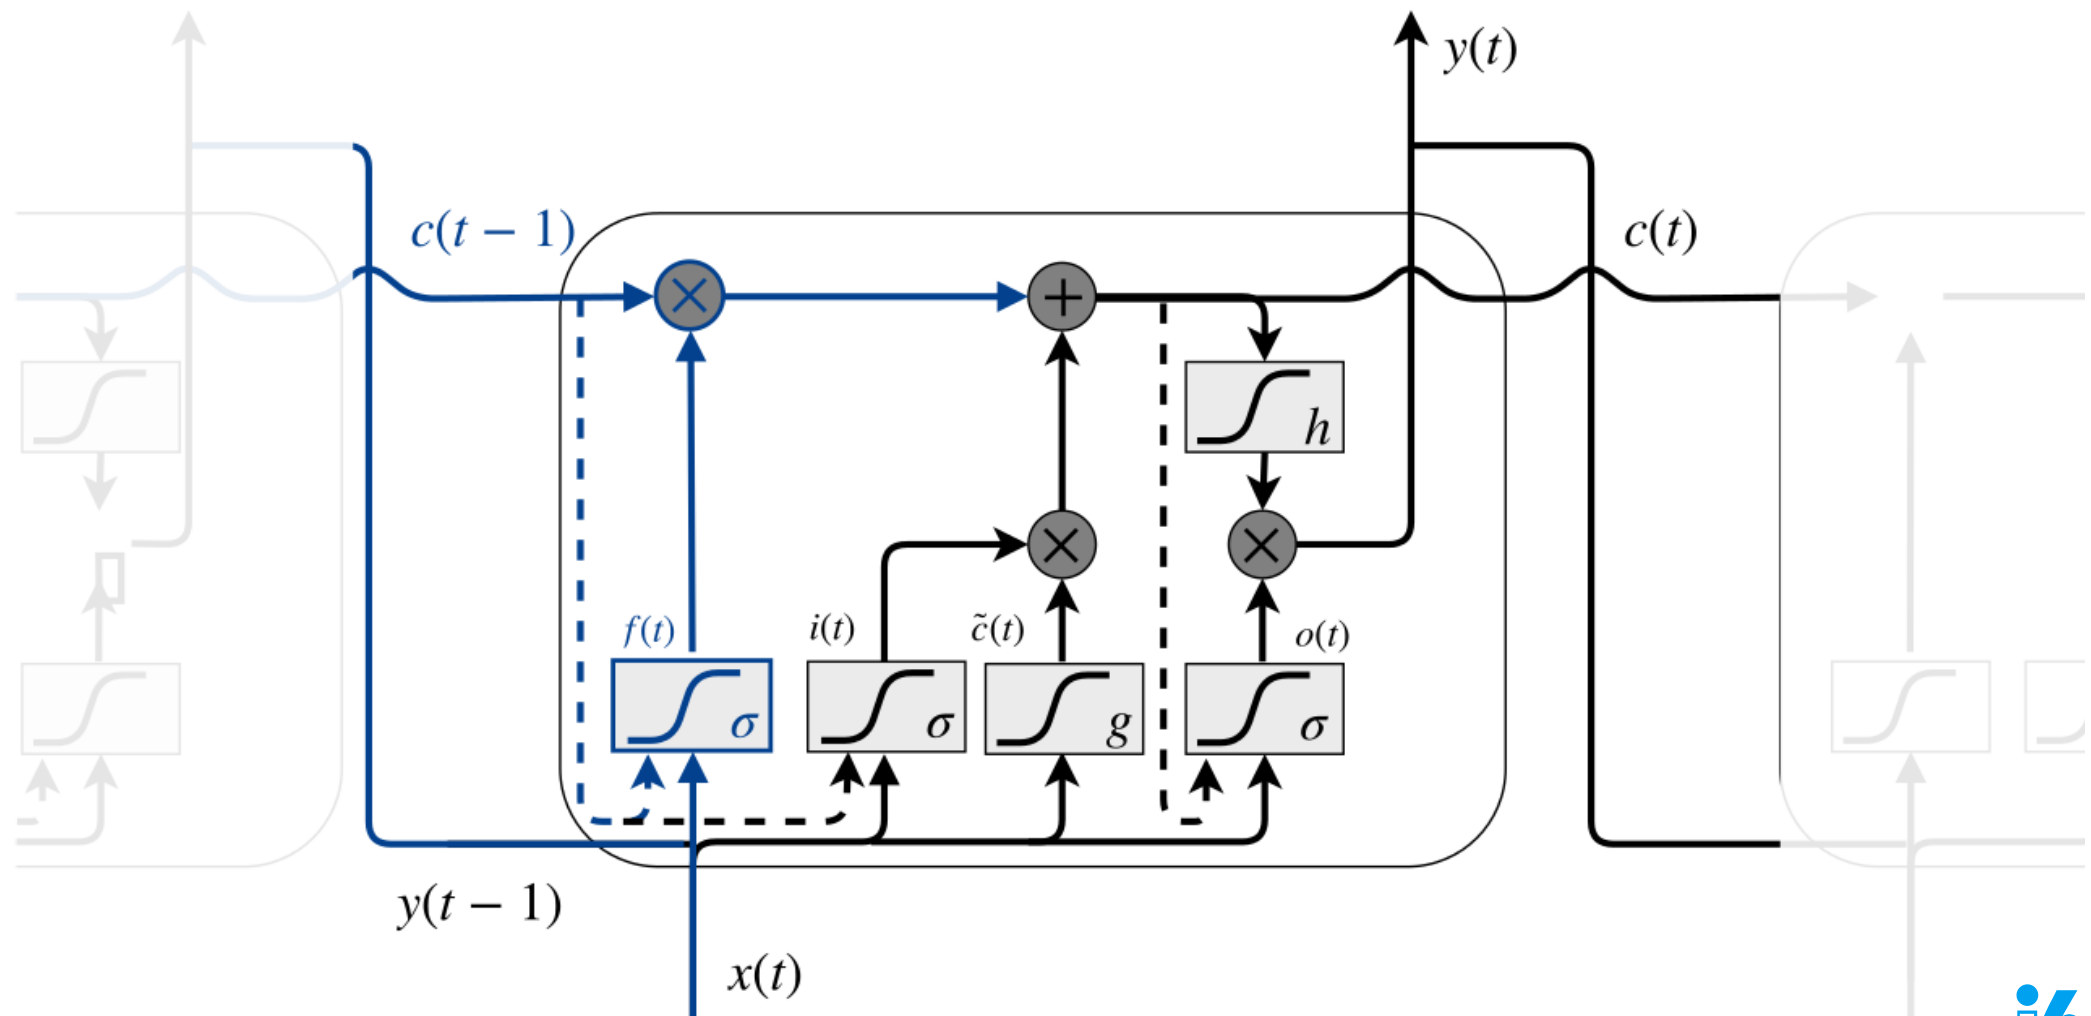
$$i(t) = \sigma(W_{xi}x(t) + W_{yi}y(t-1) + W_{ci}c(t-1) + b_i) \quad (10)$$



Long Short Term Memory (LSTM)

Forget gate:

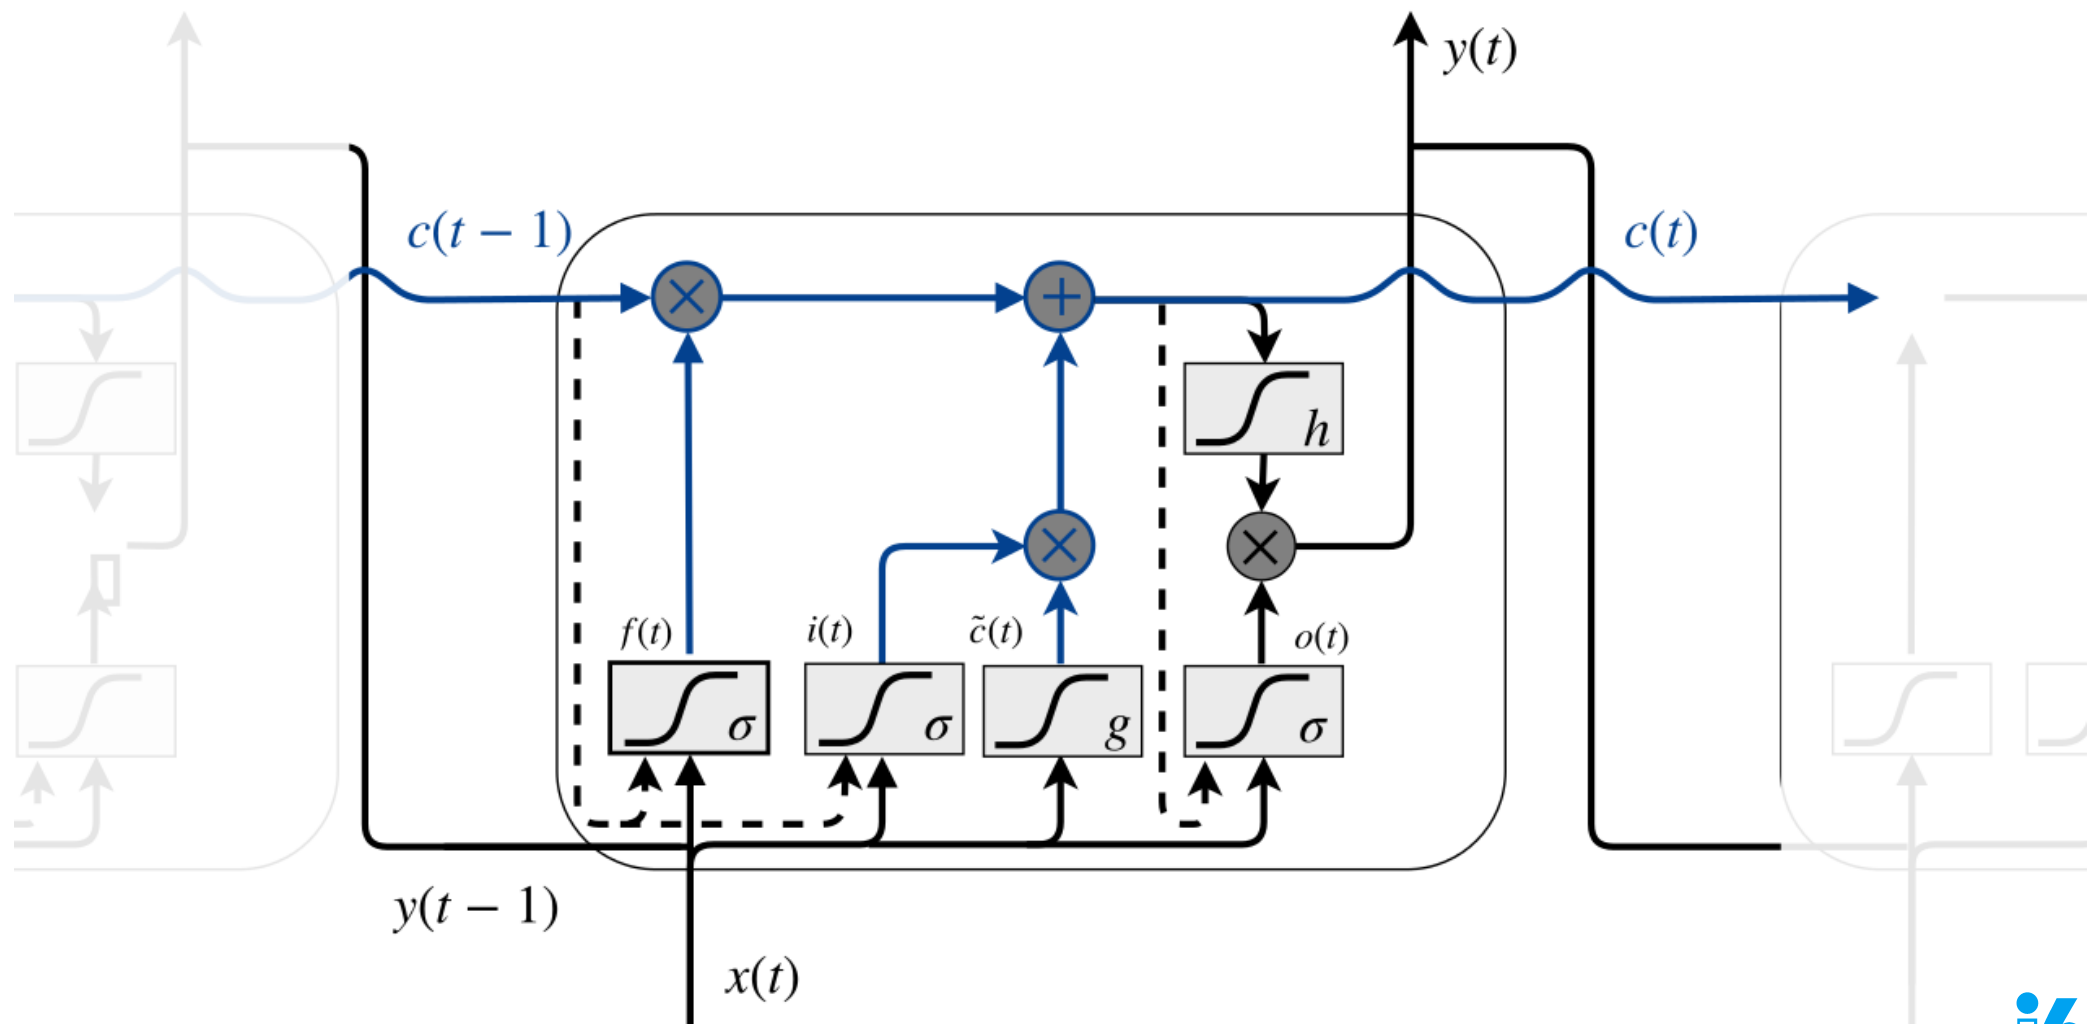
$$f(t) = \sigma(W_{xf}x(t) + W_{yf}y(t-1) + W_{cf}c(t-1) + b_f) \quad (11)$$



Long Short Term Memory (LSTM)

Updated state:

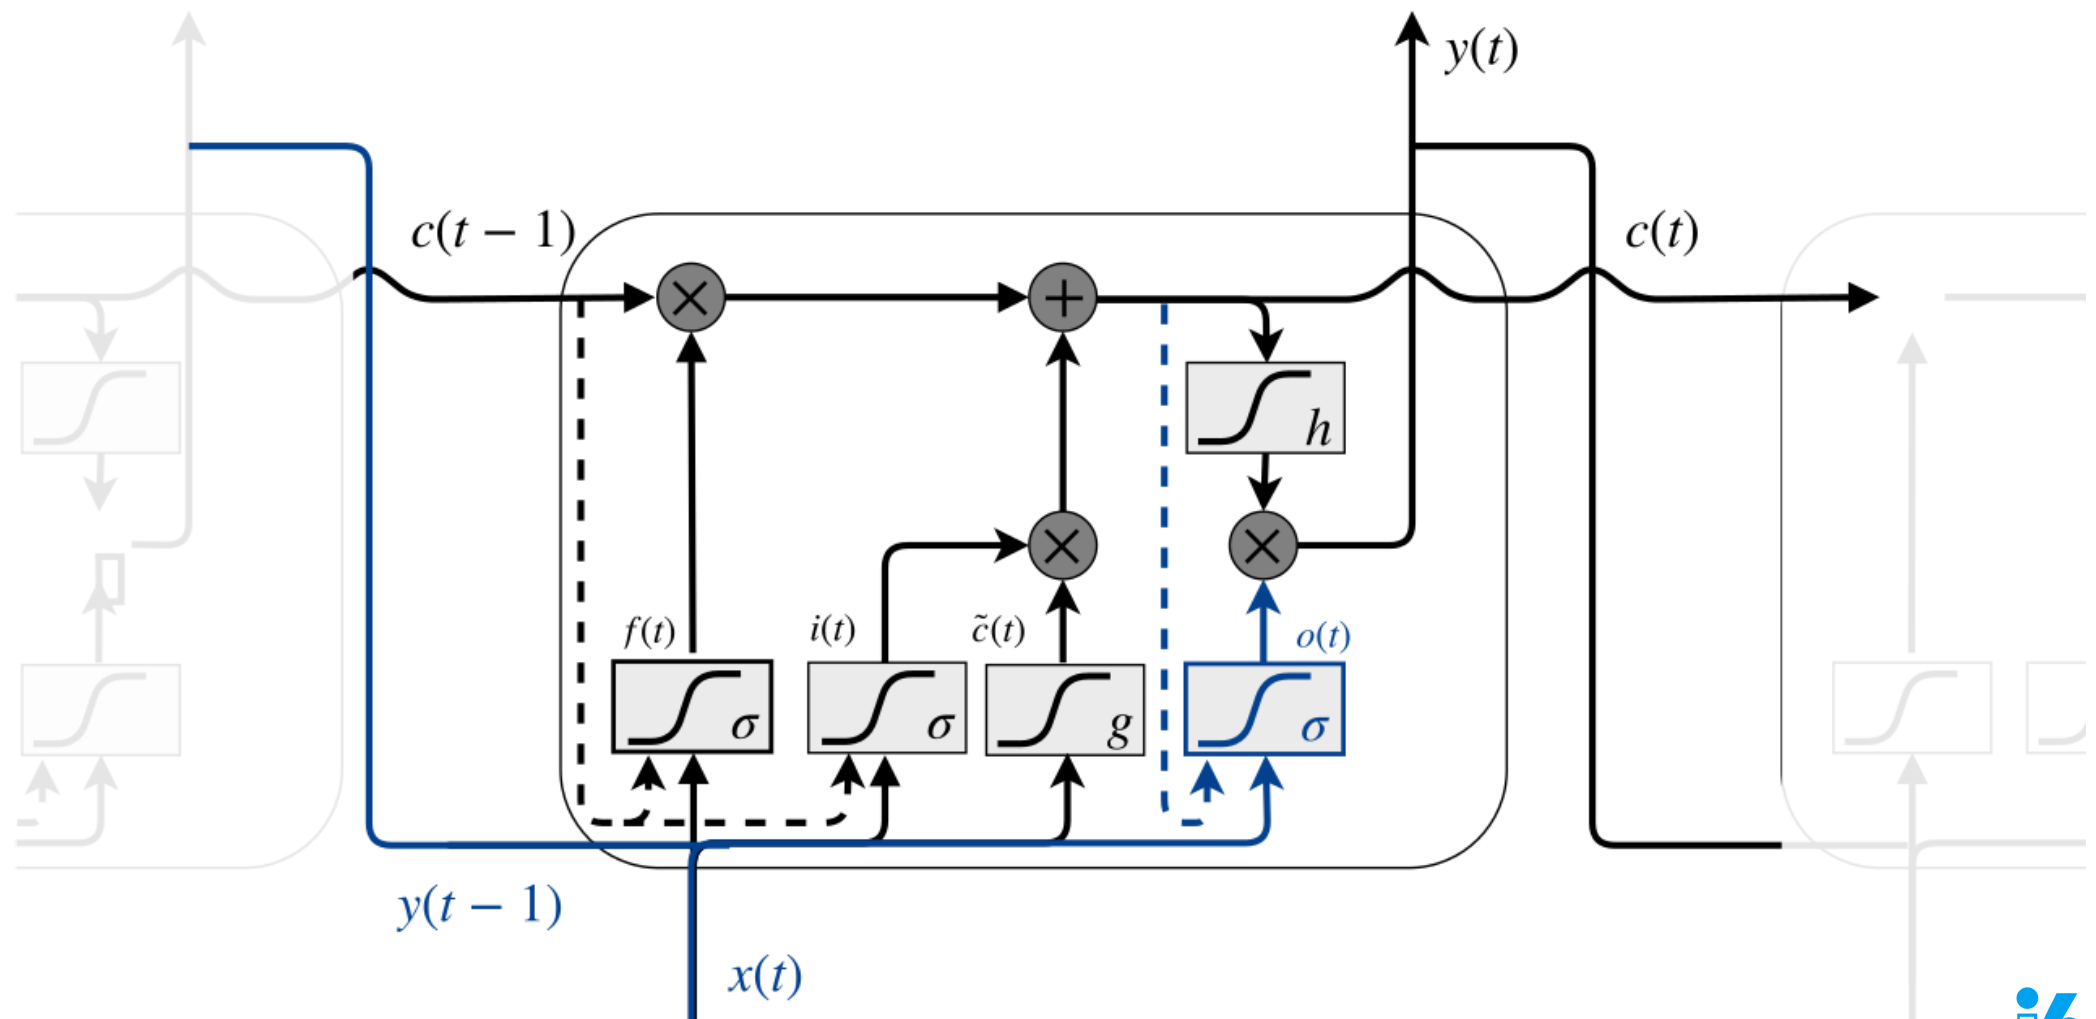
$$c(t) = f(t) \cdot c(t - 1) + i(t) \cdot \tilde{c}(t) \quad (12)$$



Long Short Term Memory (LSTM)

Output gate:

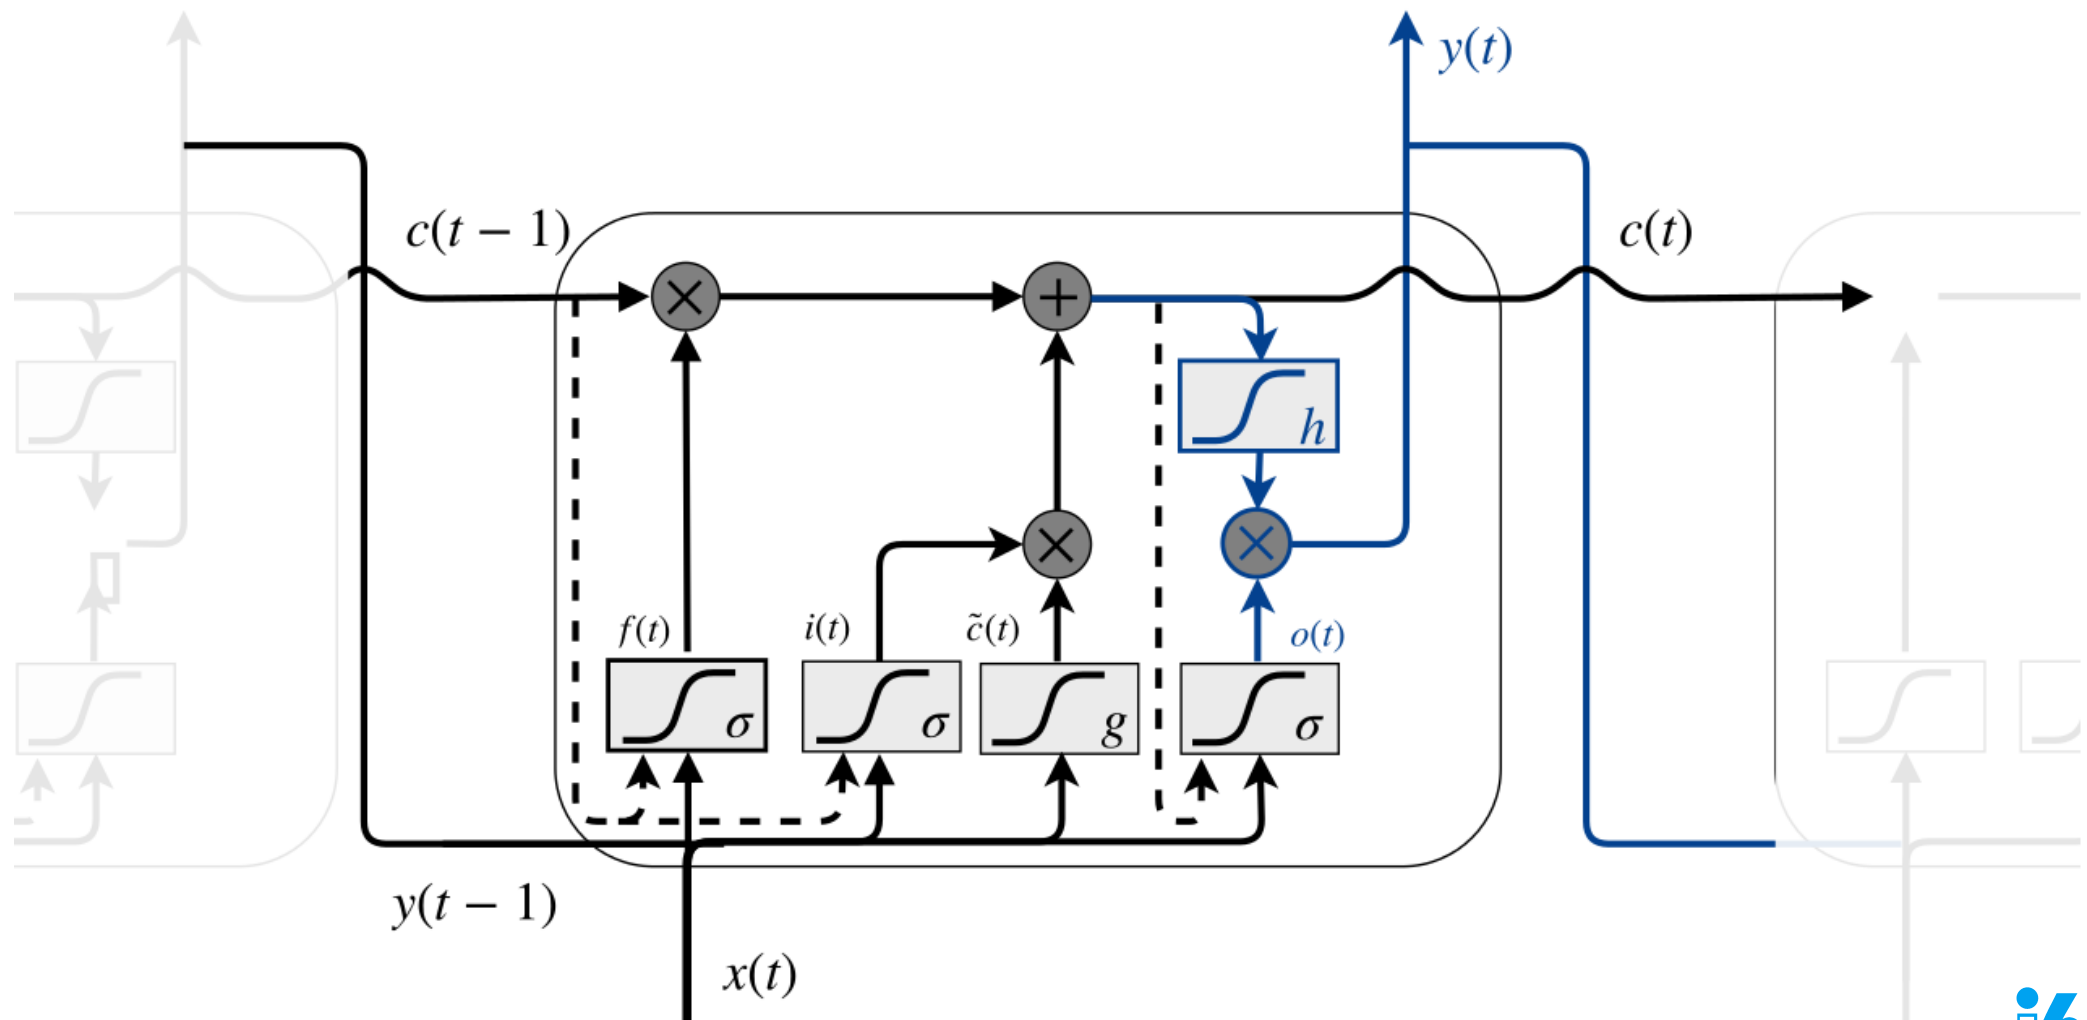
$$o(t) = \sigma(W_{xo}x(t) + W_{yo}y(t-1) + W_{co}c(t) + b_o) \quad (13)$$



Long Short Term Memory (LSTM)

Output:

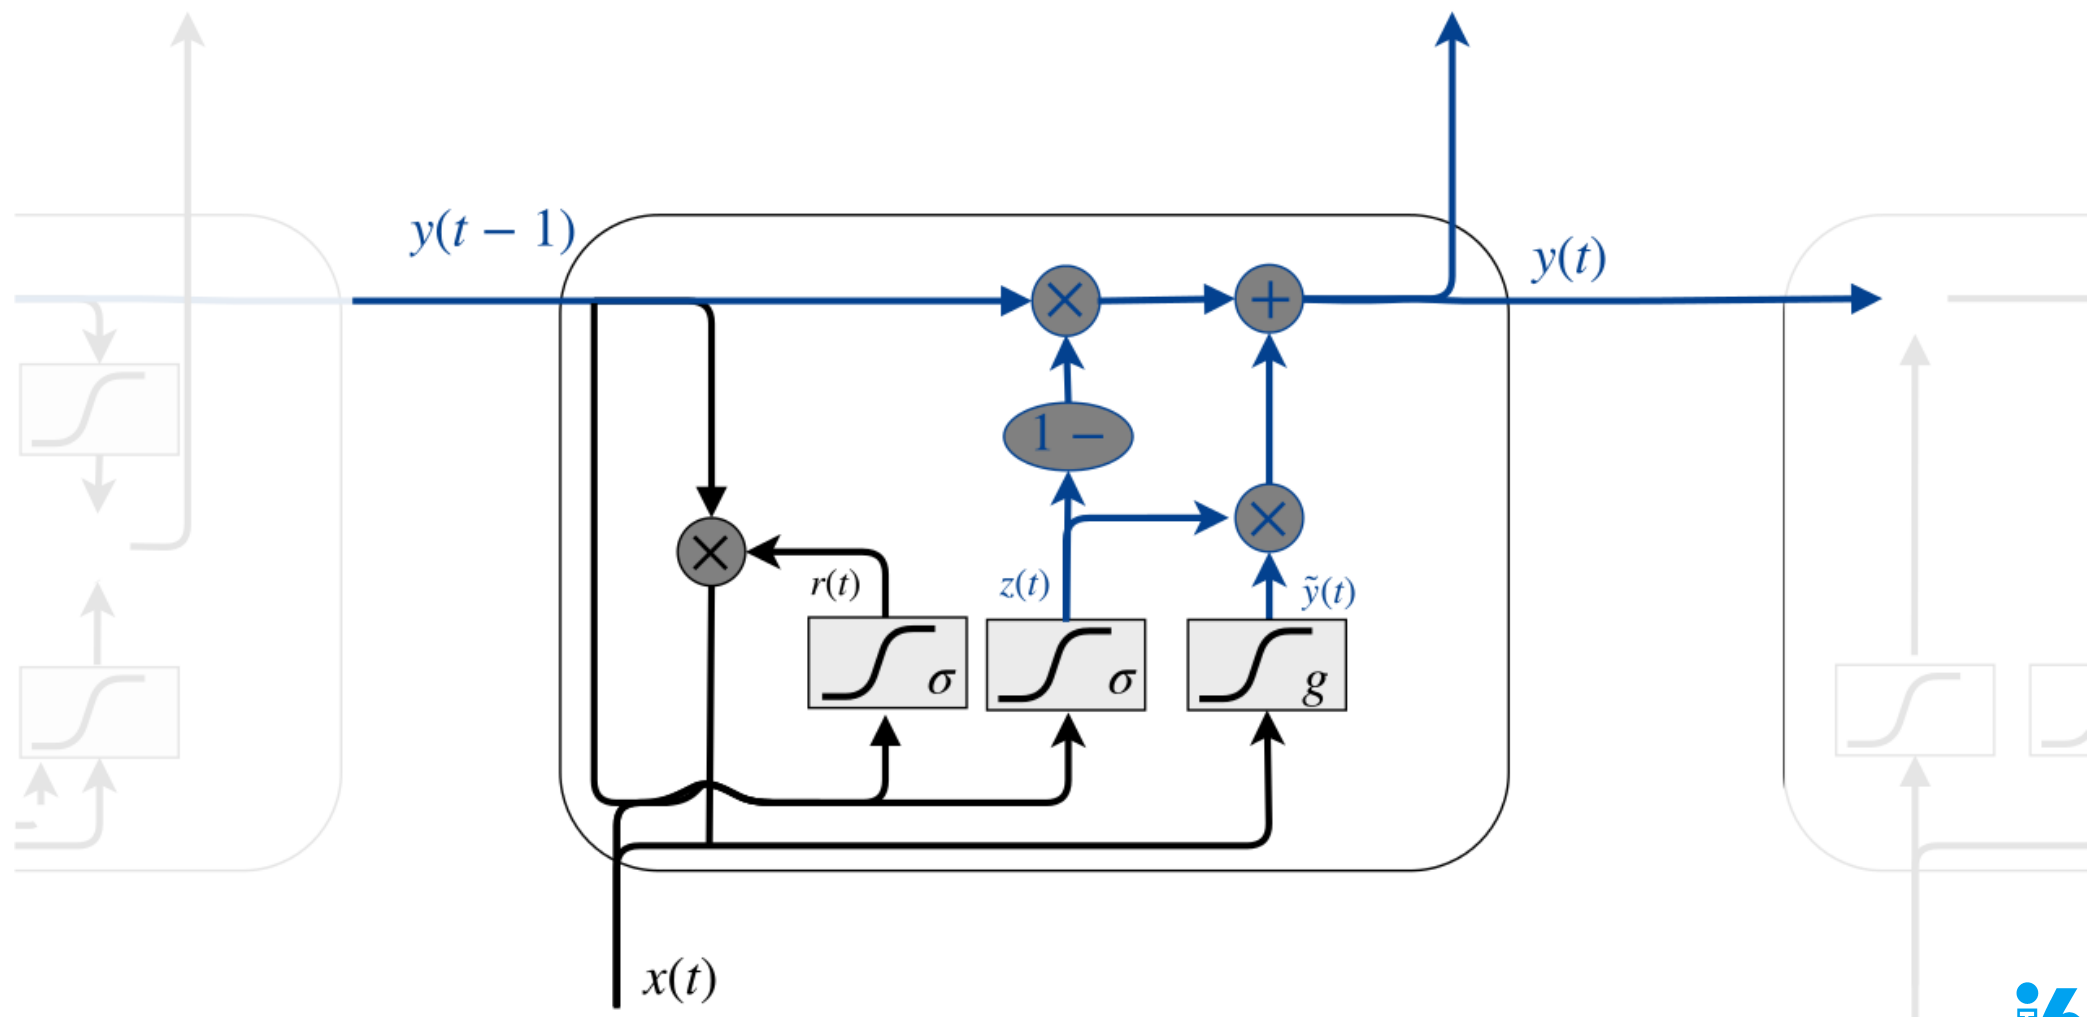
$$y(t) = o(t) \cdot h(c(t)) \quad (14)$$



Gated Recurrent Unit (GRU)

Output:

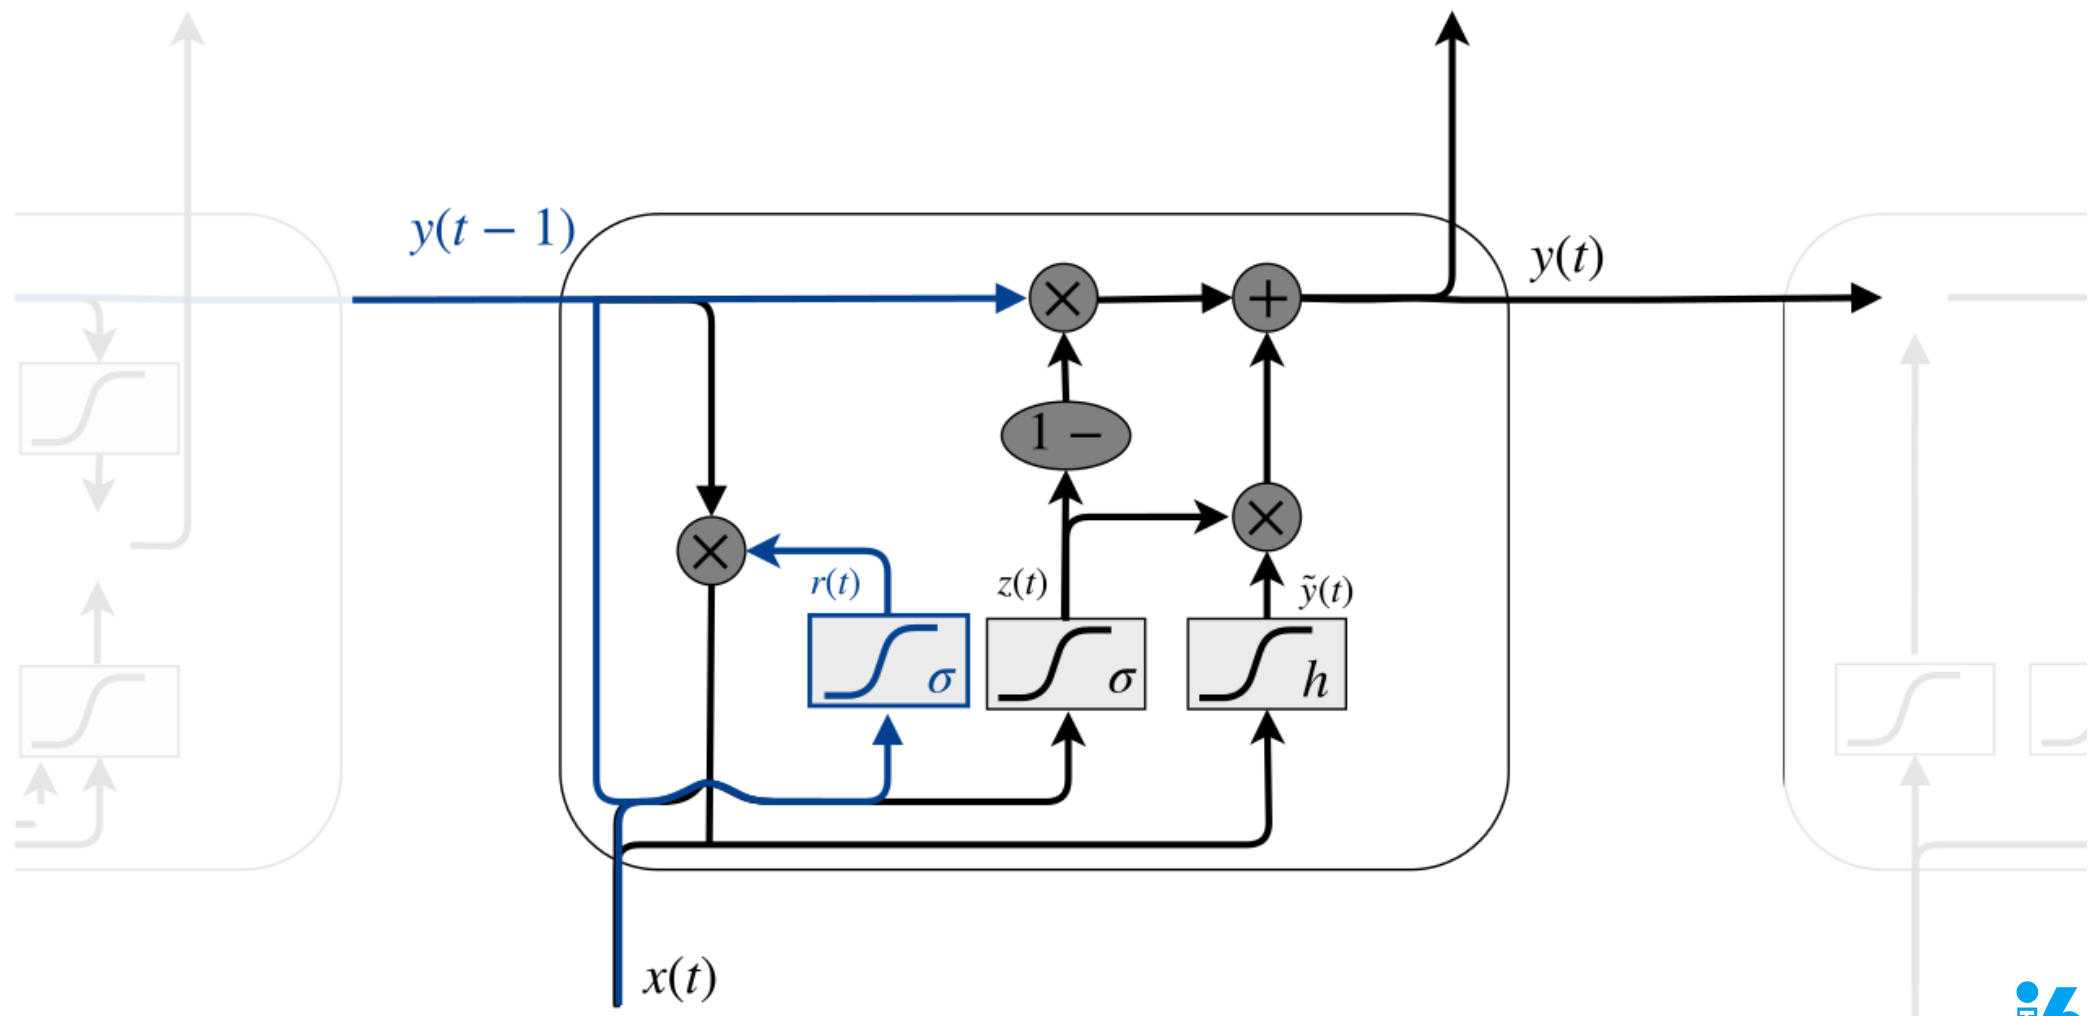
$$y(t) = (1 - z(t)) \cdot h(t - 1) + \tilde{y}(t) \cdot z(t) \quad (15)$$



Gated Recurrent Unit (GRU)

Reset gate:

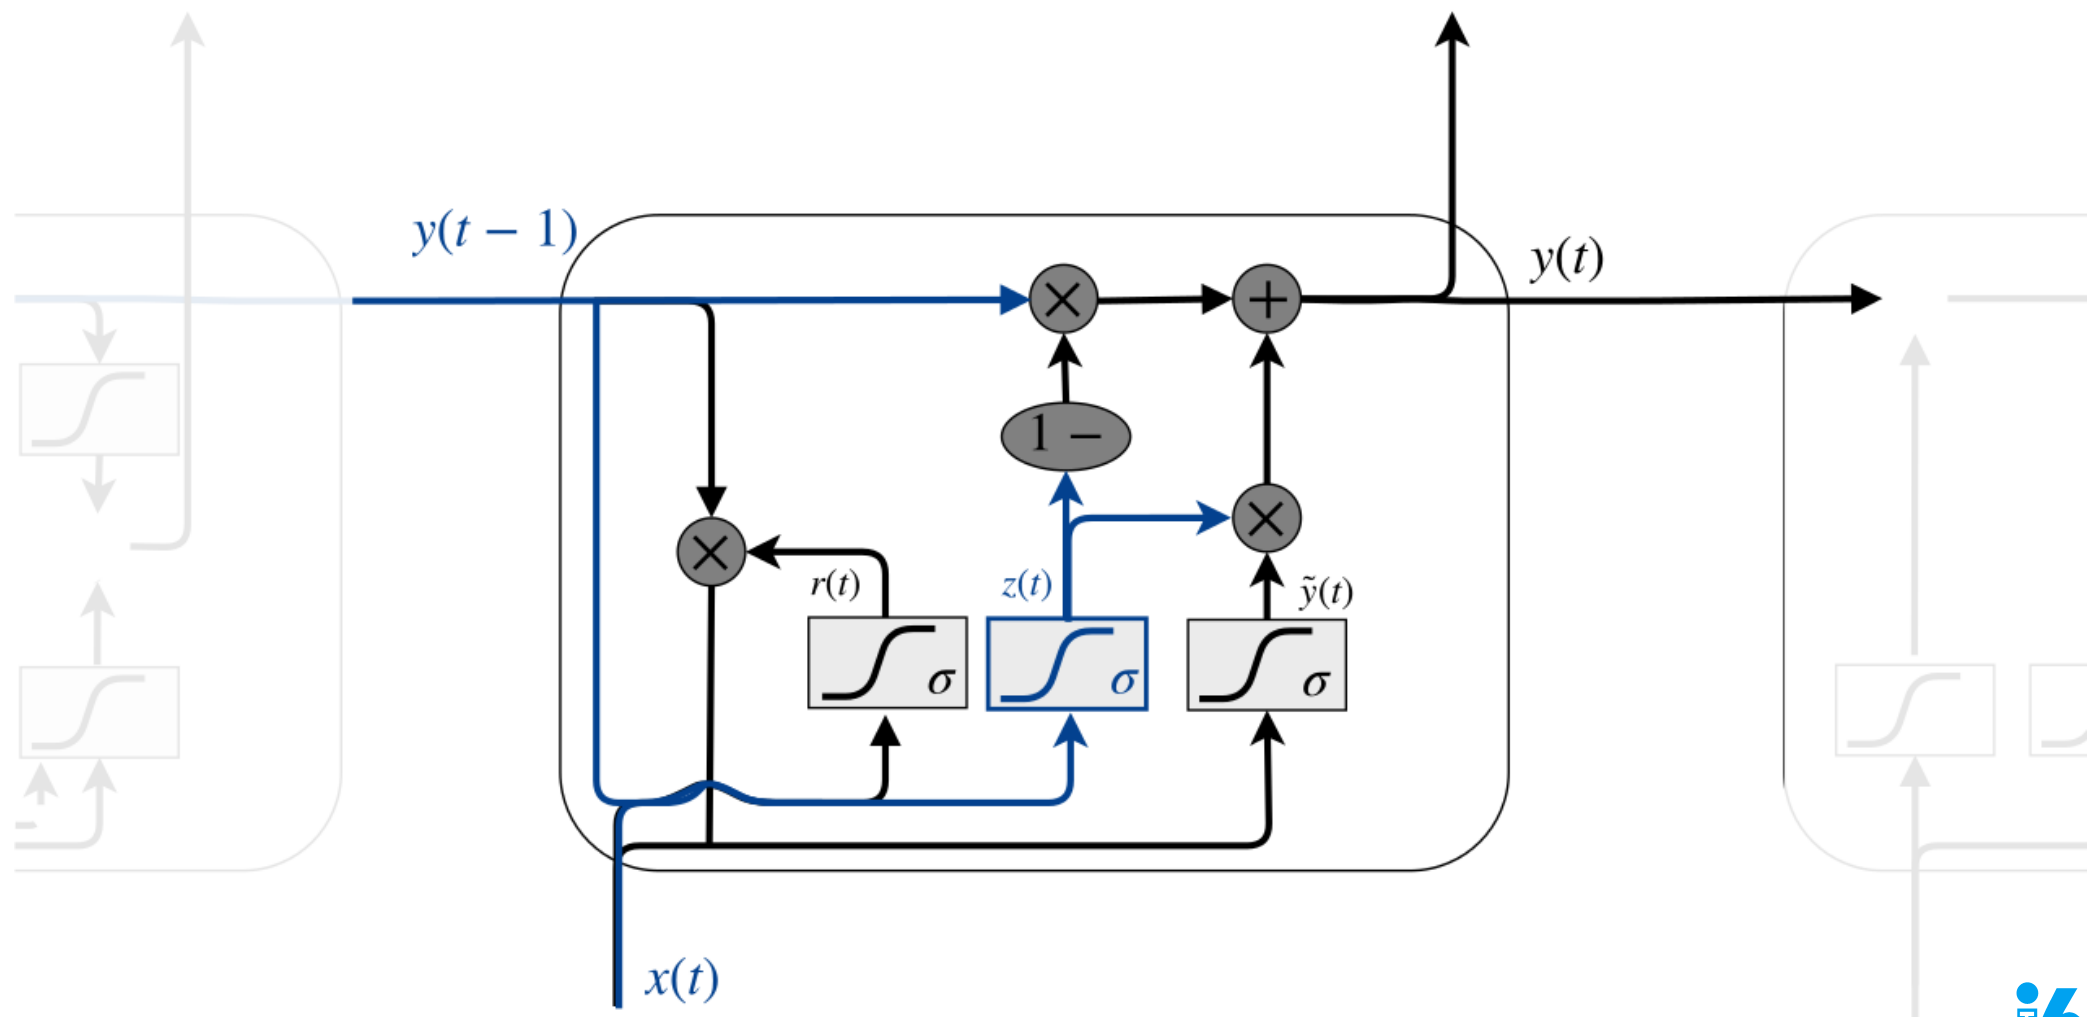
$$r(t) = \sigma(W_{xr}x(t) + W_{yr}y(t-1) + b_r) \quad (16)$$



Gated Recurrent Unit (GRU)

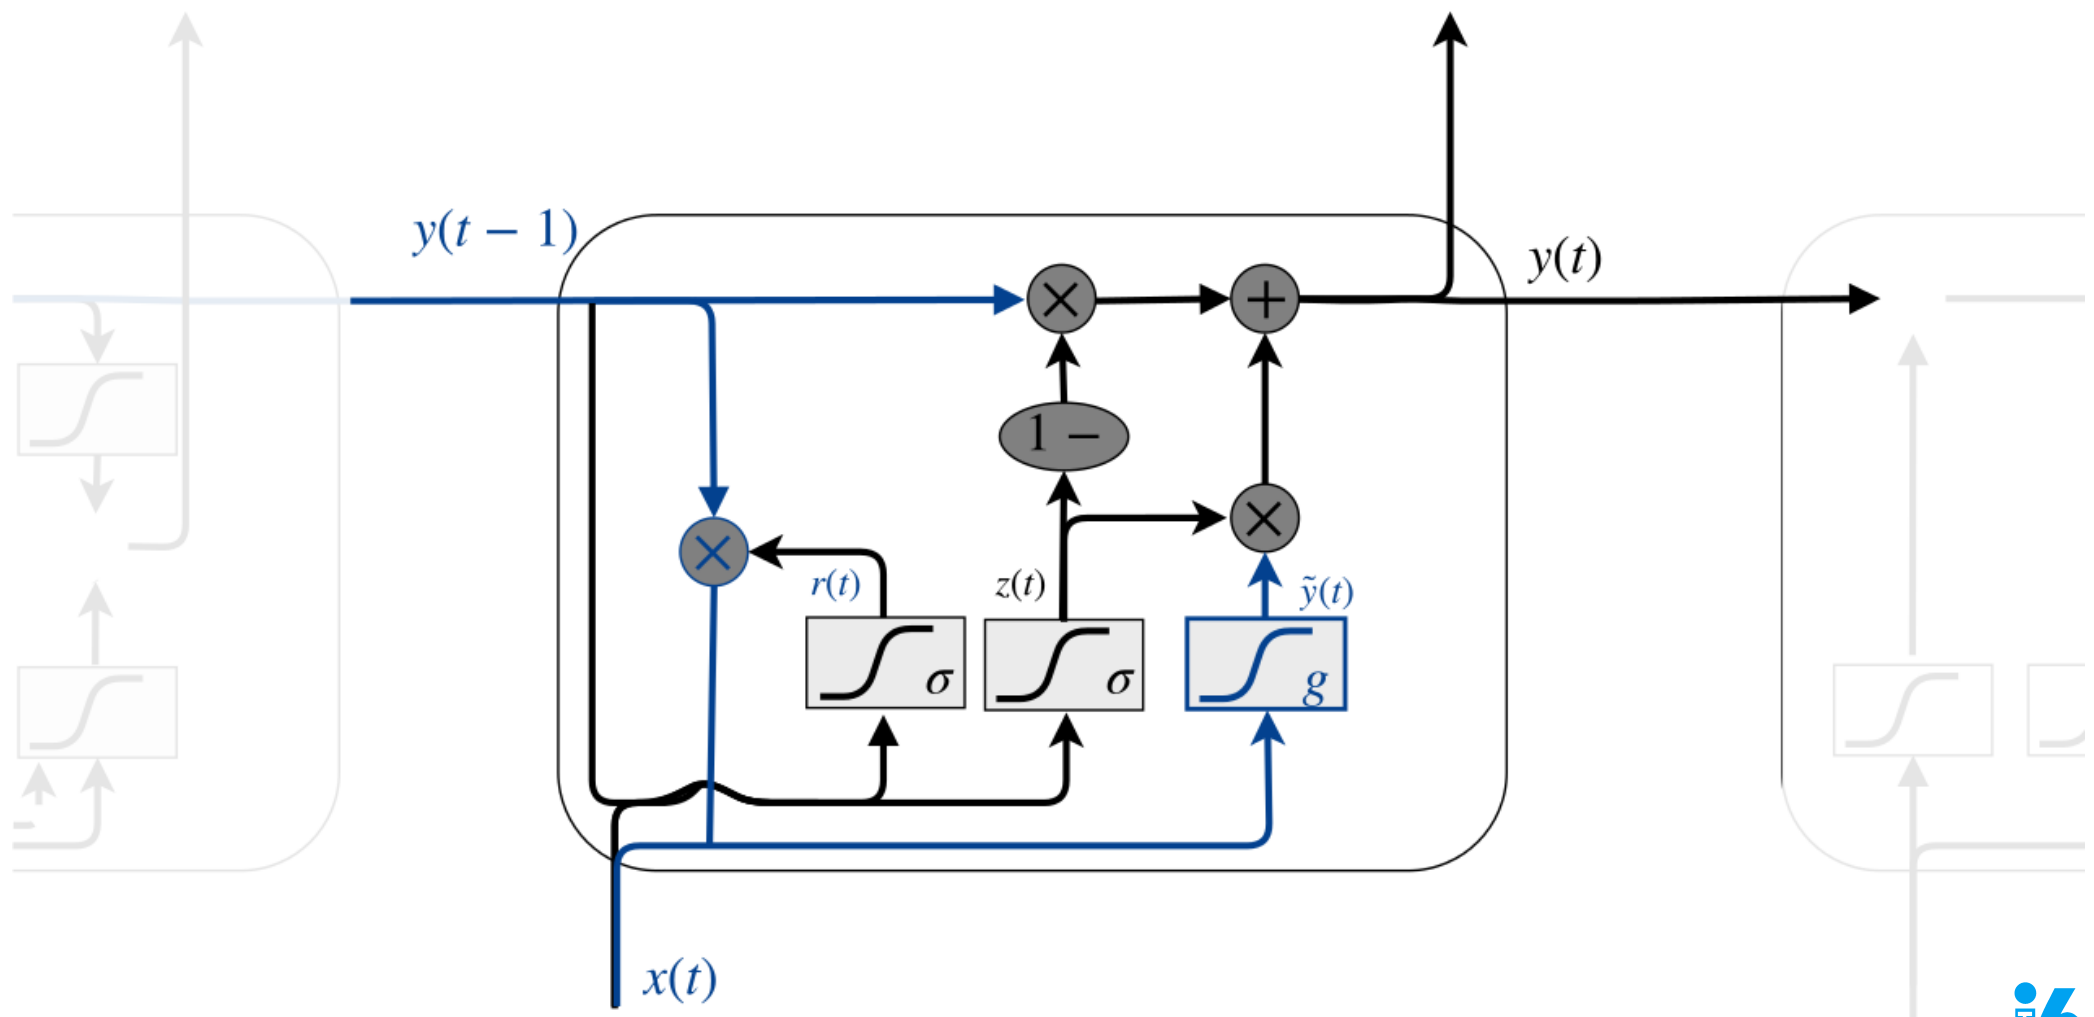
Update gate:

$$z(t) = \sigma(W_{xz}x(t) + W_{yz}y(t-1) + b_z) \quad (17)$$



Gated Recurrent Unit (GRU)

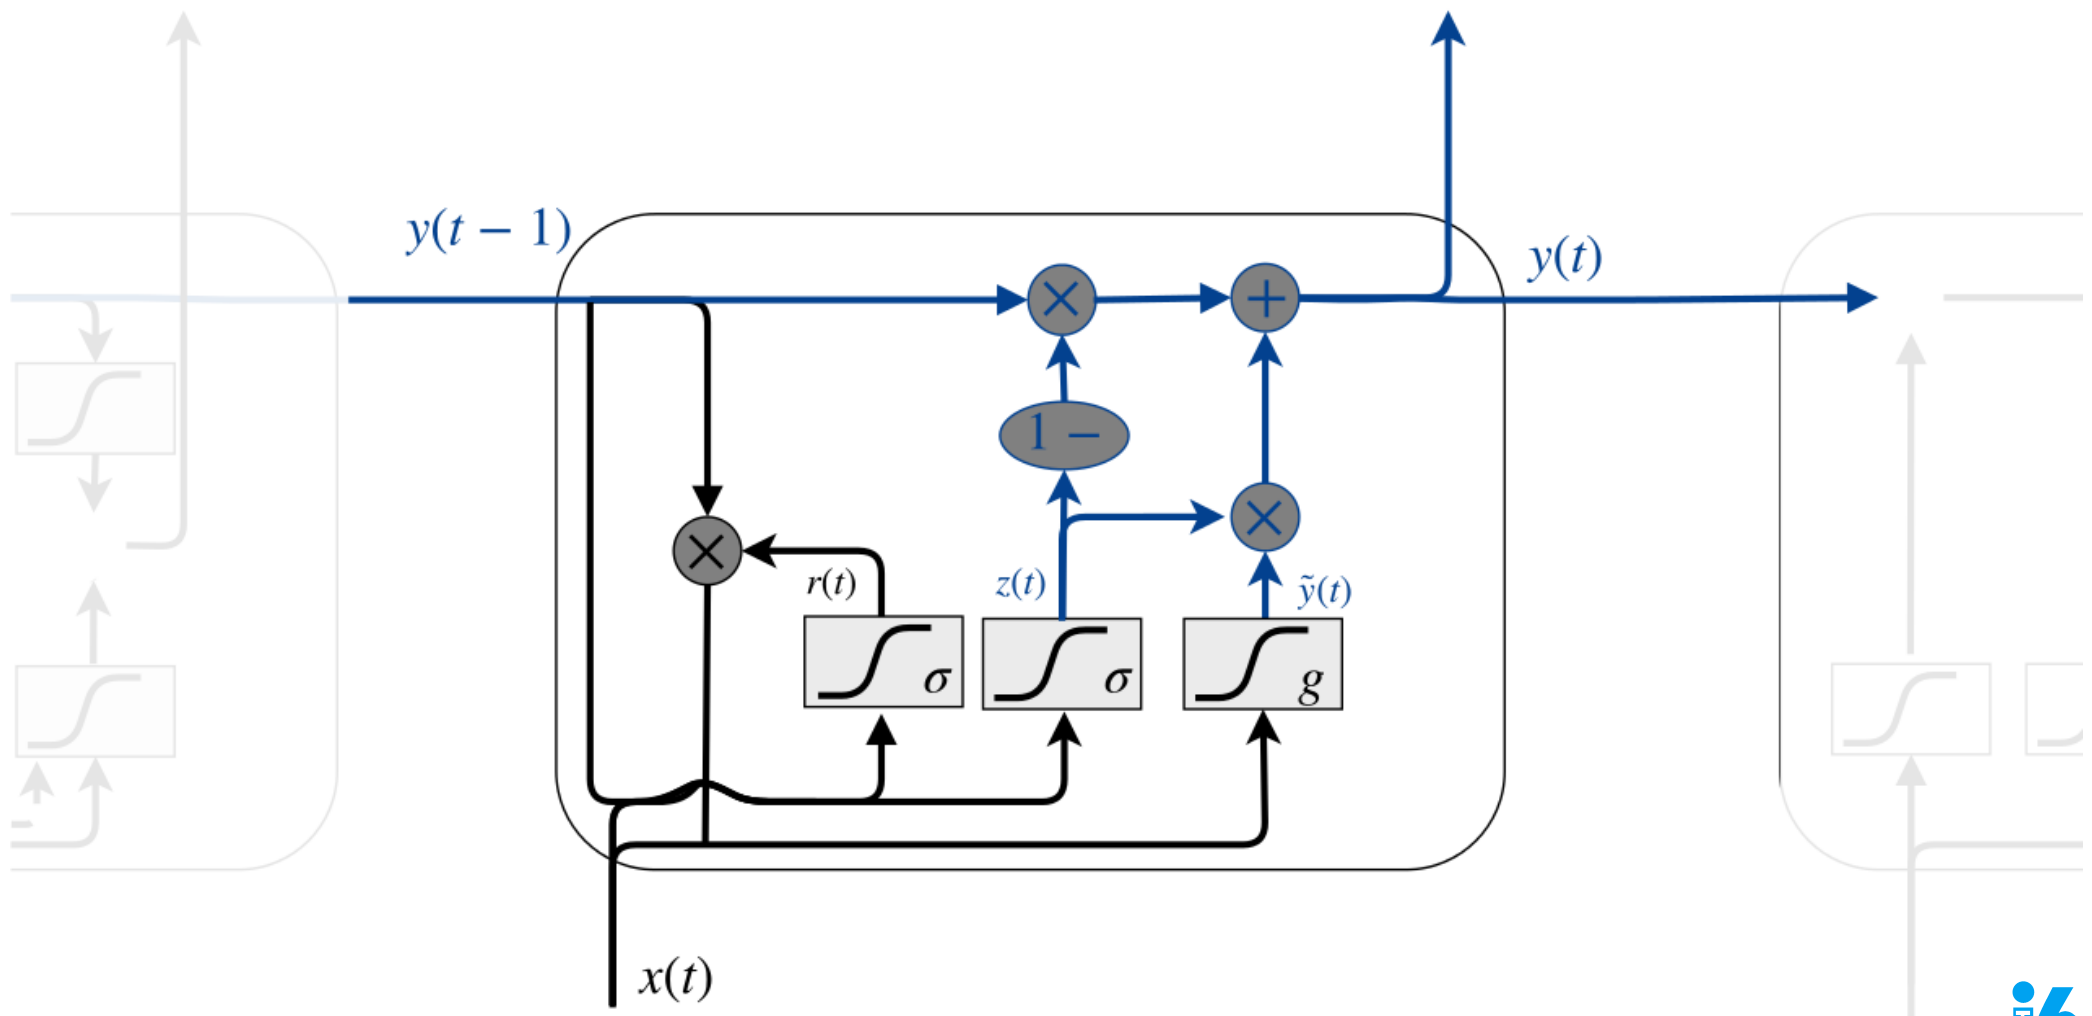
Output candidate: $\tilde{y}(t) = \tanh(W_{xy}x(t) + W_{yy}(y(t-1) \cdot r(t)) + b_y) \quad (18)$



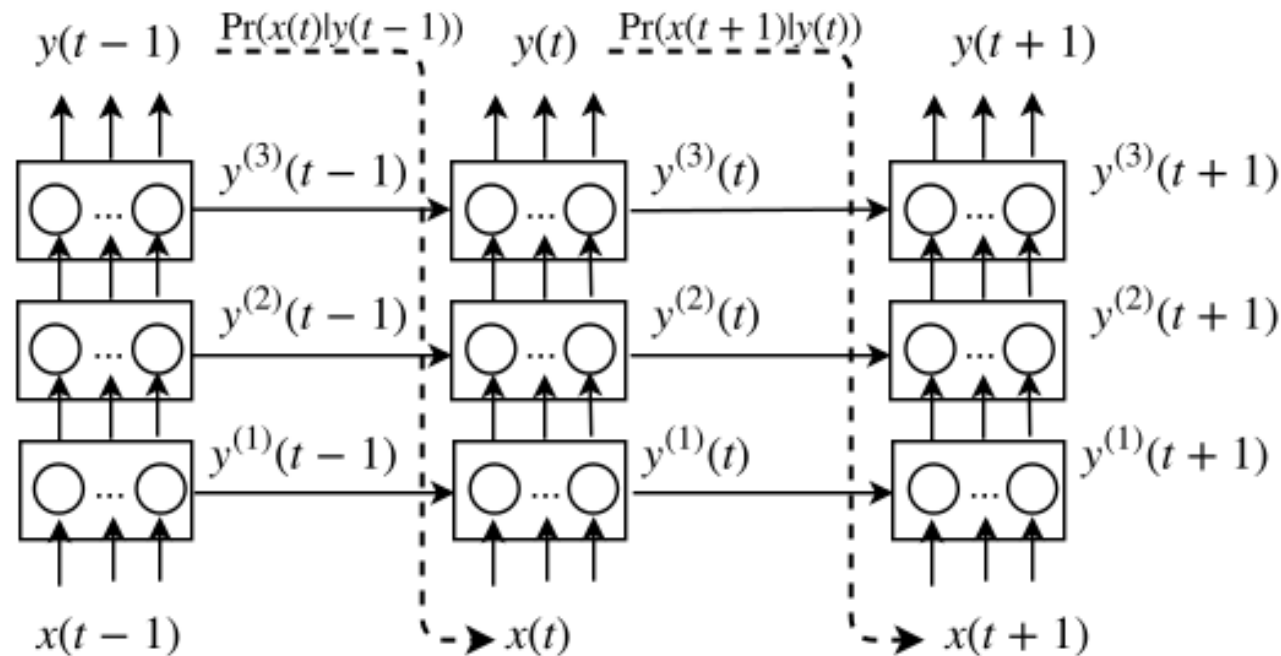
Gated Recurrent Unit (GRU)

Output:

$$y(t) = (1 - z(t)) \cdot h(t - 1) + \tilde{y}(t) \cdot z(t) \quad (19)$$



Sequence Generation [Graves 13]



- Output $y(t)$ used to predict a probability distribution of the next input $x(t+1)$

$$\mathbf{Pr}(\bar{x}) = \prod_{t=1}^T \mathbf{Pr}(x(t+1)|y(t)) \quad (20)$$

- Corresponding loss function:

$$\mathcal{L}(\bar{x}) = - \sum_{t=1}^T \log \mathbf{Pr}(x(t+1)|y(t)) \quad (21)$$

Example: Text Generation

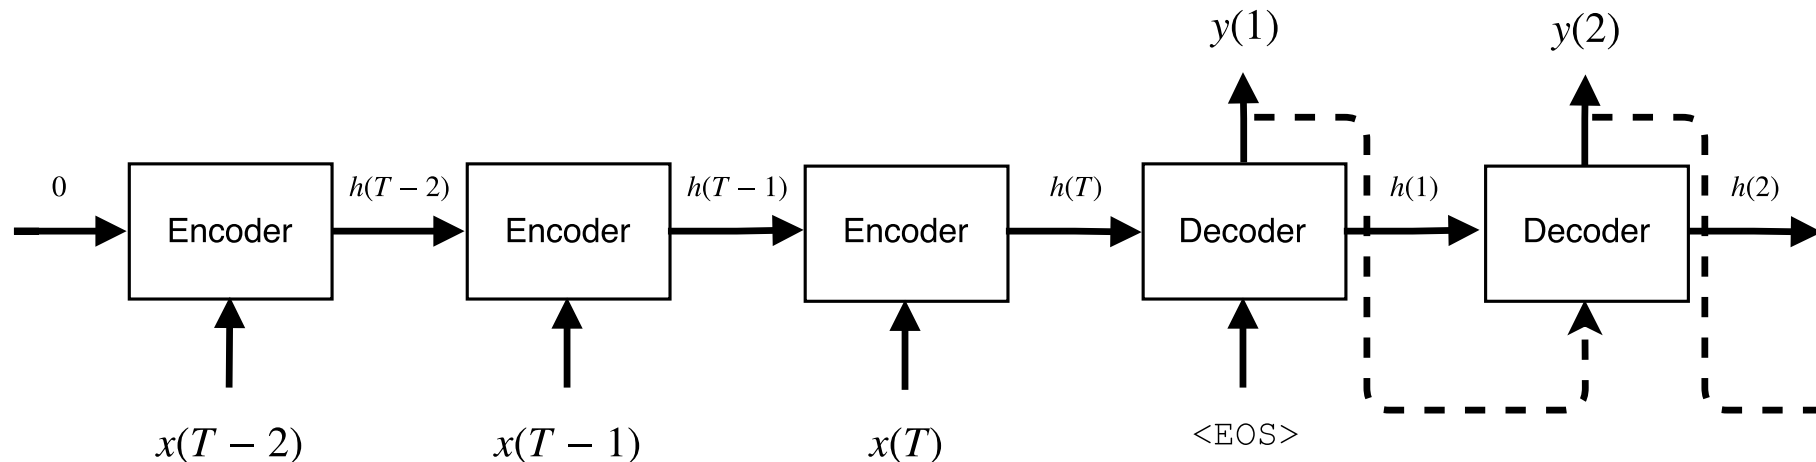
- ▶ Text represented using “one hot” encoding
 - ▷ K text classes
 - ▷ For class k : only entry k set to one and the rest to 0
- ▶ Multinomial distribution of K classes obtained by softmax layer

$$\Pr(x(t+1) = k | y(t)) = y_k(t) = \frac{\exp(\hat{y}_k(t))}{\sum_{k'=1}^K \exp(\hat{y}_{k'}(t))} \quad (22)$$

===The various disputes between Basic Mass and Council
Conditioners - "Titanist" class streams and
anarchism===

Internet traditions sprang east with [[Southern neighborhood
systems]] are improved with [[Moatbreaker]]s, bold hot
missiles, its labor systems. [[KCD]] numbered former
ISBN/MAS/speaker attacks "M3 5", which are saved as
the ballistic misely known and most functional factories.

RNN-Encoder and Decoder Model



► General model for sequence to sequence learning [Sutskever & Vinyals⁺ 14]

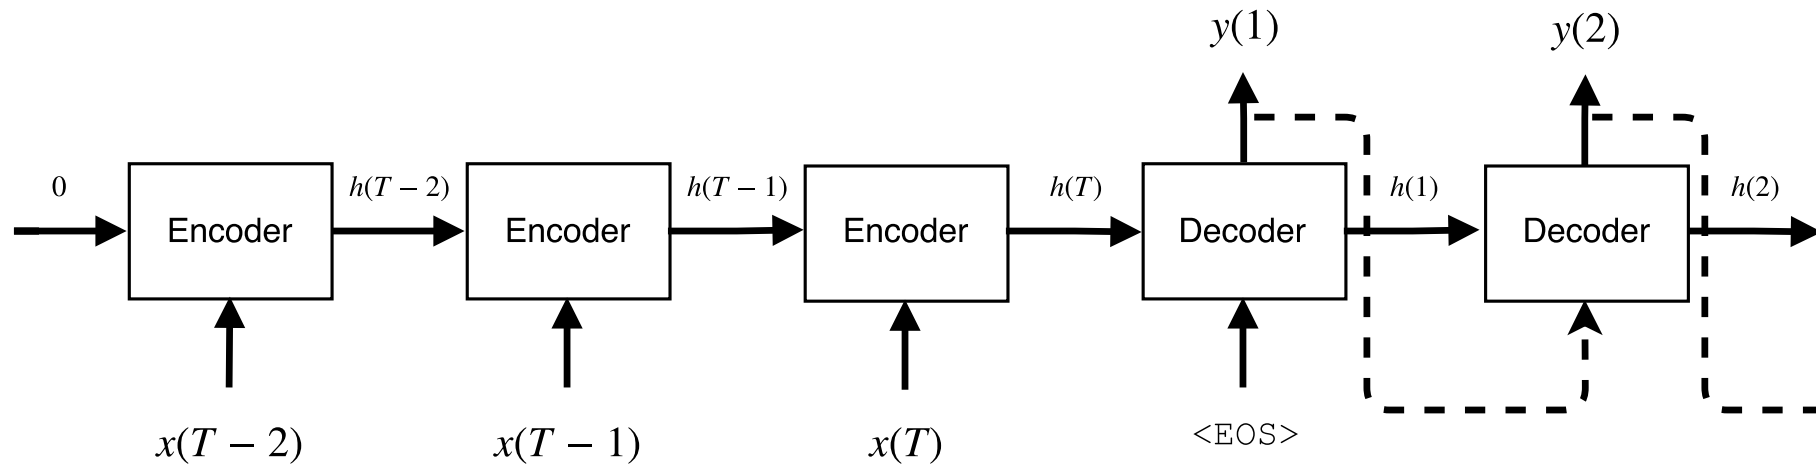
► Two separate RNNs:

- Encoder: encodes input $\bar{x} = (x(1), \dots, x(T))$ into “*thought vector*” v
- Decoder: generates output $\bar{y} = (y(1), \dots, y(T'))$ from “*thought vector*” v

$$\Pr(y(1), \dots, y(T') | x(1), \dots, x(T)) = \prod_{t=1}^{T'} \Pr(y(t) | v, y(1), \dots, y(t-1)) \quad (23)$$

► Thought vector v : fixed size representation given by $h(T)$ of encoder at T

RNN-Encoder and Decoder Model



► **Loss function:**

$$\mathcal{L}(\bar{x}, \bar{y}) = -\log \mathbf{Pr}(\bar{y}|\bar{x}) = -\sum_{t=1}^{T'} \log \mathbf{Pr}(y(t)|v, y(1), \dots, y(t-1)) \quad (24)$$

- **Encoder and decoder are LSTMs which do not share any parameters**
- **Input: word vectors**
- **Output: softmax function over all words in the vocabulary**
- **Reversed order of the input sequence leads to better performance**

Conversation Modelling

- ▶ Turn-based conversation of two parties *A* and *B*
- ▶ Task: predict the utterances of participant *B*

Objective:

- ▶ Predict the most probable utterance \bar{y}_u to follow after u utterances $\bar{x}_1, \dots, \bar{x}_u$ of participant *A* and $u - 1$ utterances $\bar{y}_1, \dots, \bar{y}_{u-1}$ of participant *B*

$$\operatorname{argmax}_{y_u} \Pr(\bar{y}_u | \bar{x}_1, \bar{y}_1, \dots, \bar{x}_{u-1}, \bar{y}_{u-1}, \bar{x}_u) \quad (25)$$

Applying the encoder-decoder framework

- ▶ Concatenate what was conversed up to point u

$$\bar{x}_u^* = \bar{x}_1 \bar{y}_1 \dots \bar{x}_{u-1} \bar{y}_{u-1} \bar{x}_u \quad (26)$$

- ▶ Use \bar{x}_u^* as input, encode it to v_u^* and generate y_u from that

$$\begin{aligned} \Pr(\bar{y}_u | \bar{x}_1, \bar{y}_1, \dots, \bar{y}_{u-1}, \bar{x}_u) &= \Pr(\bar{y}_u | \bar{x}_u^*) \\ &= \prod_{t=1}^{T'} \Pr(y_u(t) | v_u^*, y_u(1), \dots, y_u(t-1)) \end{aligned} \quad (27)$$

Conversation Modelling - Experiments [Vinyals & Le 15]

Closed-domain dataset:

- ▶ From IT helpdesk troubleshooting chats
- ▶ 30 million tokens
- ▶ Single layer LSTM with 1024 memory cells

Open-domain dataset:

- ▶ Sentences uttered by movie characters out of the *OpenSubtitles* dataset
- ▶ 923 million tokens (60 million sentences)
- ▶ Two-layered LSTM with 4096 memory cells in each layer

Evaluation:

- ▶ Comparison with rule-based chatbot (CleverBot)
- ▶ Human evaluations on a set of 200 questions

Conversation Modelling - Experiment Results

Human: what do you think about tesla ?

Machine: he 's a good conductor .

Human: what do you think about bill gates ?

Machine: he 's a good man .

Human: what is the purpose of life ?

Machine: to serve the greater good .

Human: what is the purpose of living ?

Machine: to live forever .

Human: my name is david . what is my name ?

Machine: david .

Human: how many legs does a cat have ?

Machine: four , i think .

Conversation Modelling - Experiment Results

Human: what is your job ?

Machine: i 'm a lawyer .

Human: what do you do ?

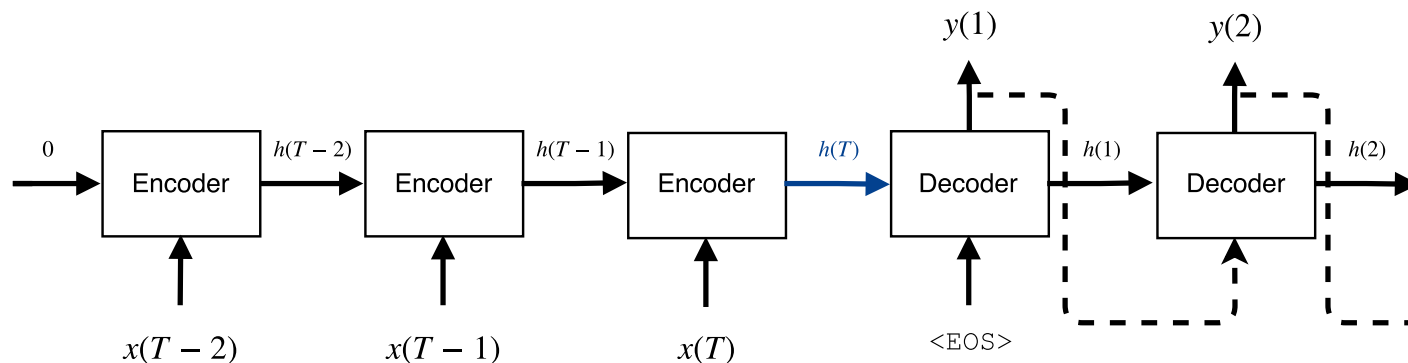
Machine: i 'm a doctor .

- ▶ Sequence to sequence model preferred in 97 out of 200 questions
- ▶ CleverBot picked in 60 out of 200
- ▶ Tie in 20 questions
- ▶ 23 questions the judges in disagreement

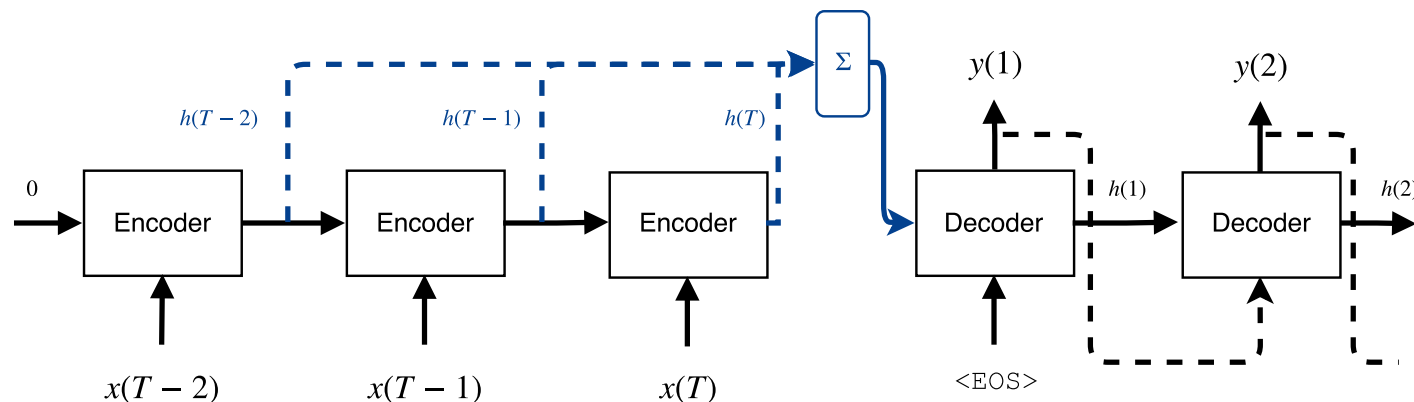
Alternative Approaches

Variants of the encoder-decoder model [Shang & Lu⁺ 15]

- **Global scheme: same approach as encoder-decoder framework**



- **Local scheme: weighted sum over all hidden states as thought vector**



- **Hybrid scheme: combines global and local scheme**

Alternative Approaches

Variants of the encoder-decoder model [Shang & Lu⁺ 15]

- ▶ **Global scheme: same approach as encoder-decoder framework**
- ▶ **Local scheme: weighted sum over all hidden states as thought vector**
- ▶ **Hybrid scheme: combines global and local scheme**
- ▶ **Experiments:**
 - ▷ **Trained on 4 million conversations from microblogging service Weibo**
 - ▷ **Implemented using GRUs instead of LSTMs**
 - ▷ **Human evaluation on 110 posts**
 - ▷ **Local scheme outperformed global scheme**
 - ▷ **Hybrid scheme beats both in all cases**

Alternative Approaches

Classification approach [Lowe & Pow⁺ 15]

- ▶ Two RNNs:
 - ▷ One encodes the context \bar{x} into fixed dimensional representation c
 - ▷ One encodes the response \bar{y} into fixed dimensional representation r
- ▶ Calculate the probability that \bar{y} is a valid response to \bar{x}

$$\Pr(\text{valid}|c, r) = \sigma(c^T M r + b) \quad (28)$$

- ▶ Trained model parameters M and b
- ▶ Can be seen as a generative approach:
 - ▷ Generate a response r s.t. $c' = M r$ is as close as possible to c

Evaluation

- ▶ **Bilingual Evaluation Understudy (BLEU) algorithm [Papineni & Roukos⁺ 02]**
 - ▷ **Compare candidate response \mathcal{C} against reference responses $\mathcal{R} \in \text{Refs}$**

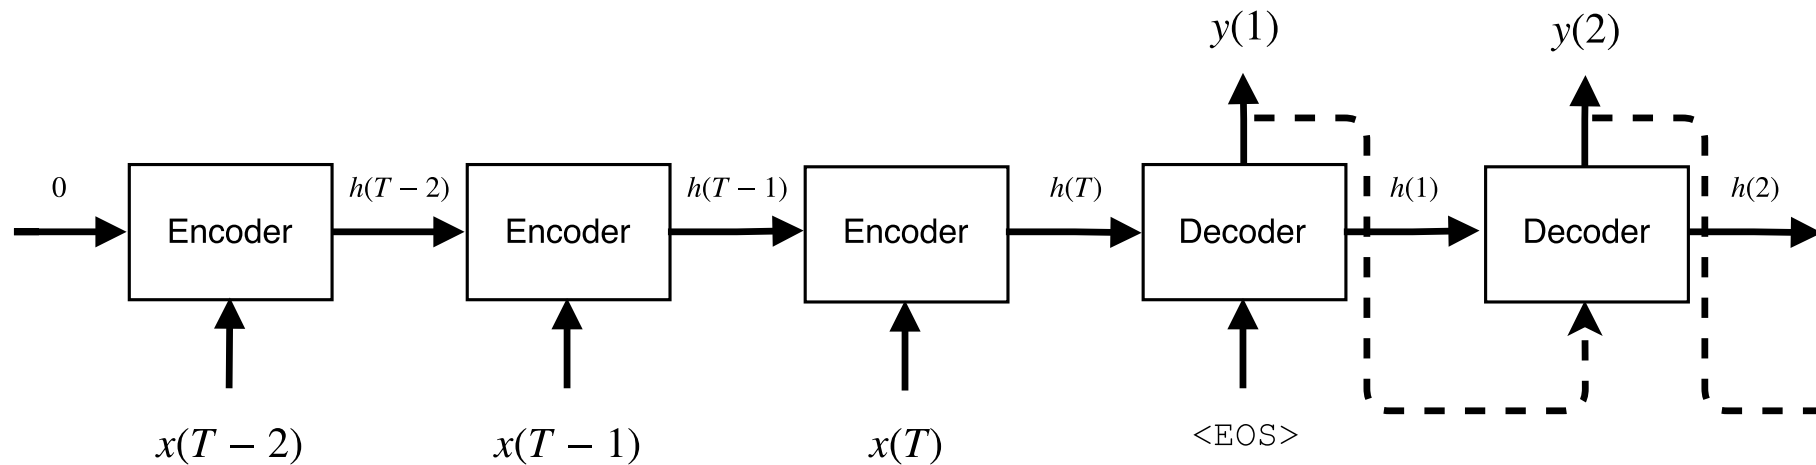
$$p_1 = \frac{\sum_{x \in \mathcal{C}} \max\{count_{\mathcal{C}}(x), \max_{\mathcal{R} \in \text{Refs}}\{count_{\mathcal{R}}(x)\}\}}{\sum_{x \in \mathcal{C}} count_{\mathcal{C}}(x)} \quad (29)$$

- ▷ **Gives very low scores for comparison of conversation utterances**
 - ▷ **But for comparison of different architectures:**
Correlates with scores produced by human experts
- ▶ **Perplexity [Brown & Pietra⁺ 92]**
 - ▷ **Indicates how well a language model predicts a corpus $X = \{x_1, \dots, x_N\}$**

$$PP(X) = \sqrt[N]{\prod_{i=1}^N \frac{1}{\Pr(x_i | x_1, \dots, x_{i-1})}} = 2^{-\frac{1}{N} \sum_{i=1}^N \log_2 \Pr(x_i | x_1, \dots, x_{i-1})} \quad (30)$$

- ▷ **Possibly not meaningful for performance in real world applications**

Conclusion



- ▶ **Encoder-decoder model suitable for various tasks of sequence generation**
- ▶ **This model adapts well to closed and open-domain datasets**
- ▶ **Hybrid and local scheme improve performance of encoder-decoder models**
- ▶ **Evaluation:**
 - ▷ **BLEU-score correlates with human judgement**
 - ▷ **Standard technique: human experts**

Thank you for your attention

Arne Nix

`arne.nix@rwth-aachen.de`

`http://arnenx.github.io/`

References

- [Brown & Pietra⁺ 92] P.F. Brown, V.J.D. Pietra, R.L. Mercer, S.A.D. Pietra, J.C. Lai: An estimate of an upper bound for the entropy of English. *Computational Linguistics*, Vol. 18, No. 1, pp. 31–40, March 1992. 22
- [Graves 13] A. Graves: Generating Sequences With Recurrent Neural Networks. *CoRR*, Vol. abs/1308.0850, August 2013. 2, 12
- [Hochreiter & Schmidhuber 97] S. Hochreiter, J. Schmidhuber: Long short-term memory. *Neural computation*, Vol. 9, No. 8, pp. 1735–1780, November 1997. 9
- [Lowe & Pow⁺ 15] R. Lowe, N. Pow, I. Serban, J. Pineau: The Ubuntu Dialogue Corpus: A Large Dataset for Research in Unstructured Multi-Turn Dialogue Systems. In *Proceedings of the SIGDIAL Conference*, pp. 285–294. Association for Computational Linguistics, September 2015. 21, 29
- [Papineni & Roukos⁺ 02] K. Papineni, S. Roukos, T. Ward, W.J. Zhu: BLEU: a method for automatic evaluation of machine translation. In *Proceedings of the 40th annual meeting on association for computational linguistics*, pp. 311–318. Association for Computational Linguistics, July 2002. 22

- [Shang & Lu⁺ 15] L. Shang, Z. Lu, H. Li: Neural Responding Machine for Short-Text Conversation. In *Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics and the 7th International Joint Conference on Natural Language Processing*, pp. 1577–1586. Association for Computational Linguistics, July 2015. 20
- [Sutskever & Vinyals⁺ 14] I. Sutskever, O. Vinyals, Q.V. Le: Sequence to Sequence Learning with Neural Networks. In *Advances in Neural Information Processing Systems 27: Annual Conference on Neural Information Processing Systems 2014, December 8-13 2014, Montreal, Quebec, Canada*, pp. 3104–3112, 2014. 2, 14
- [Vinyals & Le 15] O. Vinyals, Q.V. Le: A Neural Conversational Model. In *Proceedings of the 31st International Conference on Machine Learning*, Vol. 37, July 2015. 2, 17

Appendix: Training Neural Networks

Back Propagation Algorithm

- Recursion equation for the derivative with respect to the output of layer l :

$$\delta_i^{(l)} = \frac{\partial \mathcal{L}}{\partial y_i^{(l)}} = \sum_{j=1}^J \frac{\partial \mathcal{L}}{\partial y_j^{(l+1)}} \frac{\partial y_j^{(l+1)}}{\partial y_i^{(l)}} = \sum_{j=1}^J \delta_j^{(l+1)} \frac{\partial y_j^{(l+1)}}{\partial y_i^{(l)}} \quad (31)$$

- Derivative with respect to the weights of layer l :

$$\frac{\partial \mathcal{L}}{\partial w_{ji}^{(l)}} = \sum_{j=1}^J \frac{\partial \mathcal{L}}{\partial y_j^{(l+1)}} \frac{\partial y_j^{(l+1)}}{\partial w_{ji}^{(l)}} = \sum_{j=1}^J \delta_j^{(l+1)} \frac{\partial y_j^{(l+1)}}{\partial w_{ji}^{(l)}} \quad (32)$$

Appendix: Training Recurrent Neural Networks

Back Propagation Through Time

- Recursion equation for the derivative with respect to the output of layer l :

$$\delta_i^{(l)}(t) = \sum_{j=1}^J \delta_j^{(l+1)}(t) \frac{\partial y_j^{(l+1)}(t)}{\partial y_i^{(l)}(t)} + \sum_{k=1}^K \delta_k^{(l+1)}(t+1) \frac{\partial y_k^{(l+1)}(t+1)}{\partial y_i^{(l)}(t+1)} \quad (33)$$

- Derivative with respect to the weights of layer l :

$$\frac{\partial \mathcal{L}}{\partial w_{ji}^{(l)}} = \sum_{t=1}^T \sum_{j=1}^J \delta_j^{(l+1)}(t) \frac{\partial y_j^{(l+1)}(t)}{\partial w_{ji}^{(l)}(t)} \quad (34)$$

Appendix: Evaluation

- ▶ **Recall@k metric [Lowe & Pow⁺ 15]**
 - ▷ Model names the k most likely responses to given context
 - ▷ Output is correct if the true response is among these k
 - ▷ Not useful for generative approaches
 - ▷ Used to set a benchmark on the *Ubuntu Dialogue Corpus*
- ▶ **Human experts**
 - ▷ Ranking of responses from 0 (unsuitable) to +2 (suitable)
 - ▷ Performance criteria: grammar and fluency, logic consistency, semantic relevance, scenario dependence and generality
- ▶ **In the future**
 - ▷ Use a metric to compare generated responses with the true response
 - ▷ Dependent on a standardized embedding
 - possible candidate: *skipped thought vectors*

